

<https://gitter.im/ScalaTaiwan/ScalaTaiwan>

Scala & Spark(1.6) in Performance Aspect

Jimin Hsieh
<https://tw.linkedin.com/in/jiminhsieh>

2016/06/14 @ [Scala Taiwan Meetup](#)

Who am I?

- ❖ Server-side engineer who develops anything except UI and loves performance tuning, FP, OOP, and Linux.
- ❖ Recently, I am doing data processing with Spark Scala.
- ❖ Experience:
 - ❖ QA (Computer Network)
 - ❖ Embedded Application Engineer (VoIP)
 - ❖ Server-side Engineer

Agenda

- ❖ How to Performance tuning?
- ❖ Scala Performance
- ❖ Why Scala? (*)
- ❖ Spark(1.6) Scala Performance

How to performance tuning?

- ❖ The only way to become a better programmer is **by not programming**. Code is important, but it's a small part of the overall process. To truly become a better programmer, **you have to cultivate passion for everything else that goes on around the programming.**
From [How To Become a Better Programmer by Not Programming](#) by Jeff Atwood, co-founder of Stack Overflow
- ❖ Experiment + **Know it through and through + Patience**

Scala Performance

- ❖ Collection Performance
 - ❖ immutable(Persistent data structure) vs mutable
- ❖ Loop vs Recursion vs Combinators
- ❖ Primitive type vs Boxed primitive
- ❖ String size

Collection Performance

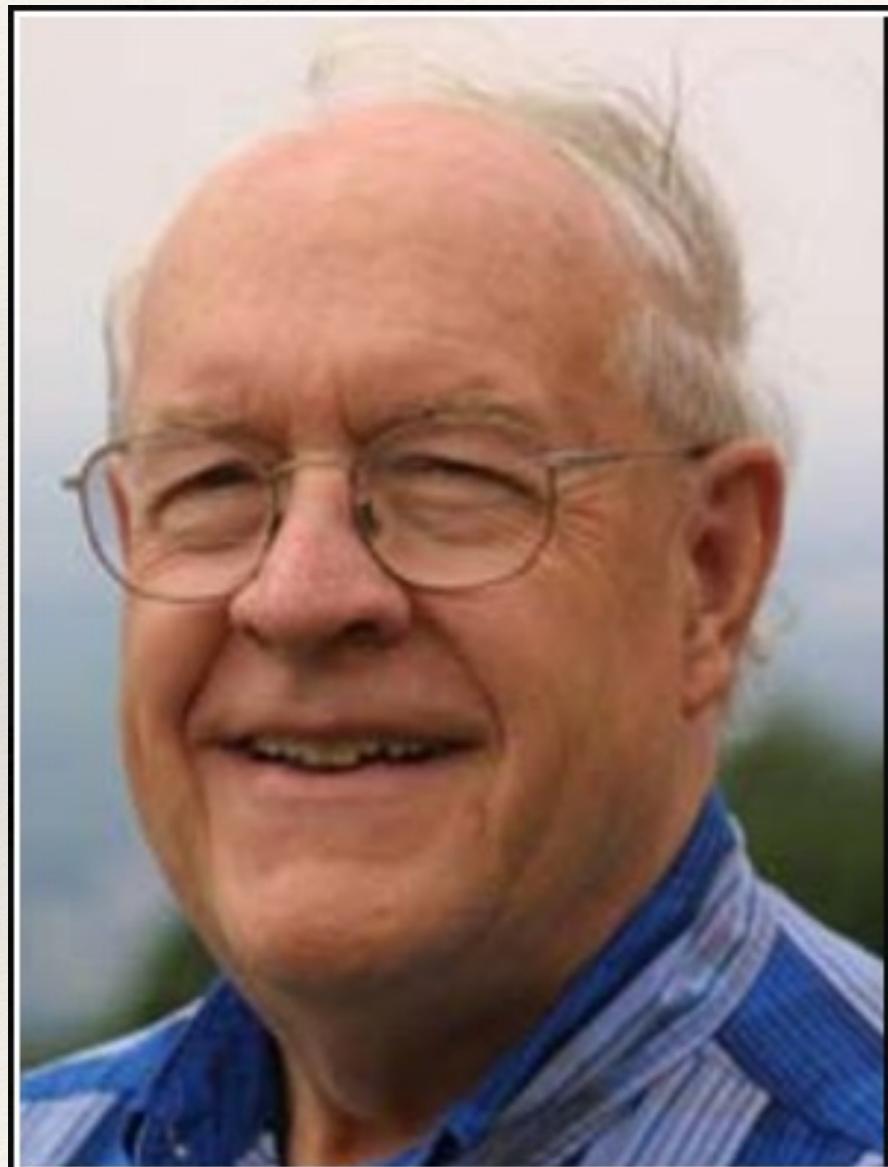


Bad programmers worry about the code. Good programmers worry about data structures and their relationships.

— *Linus Torvalds* —

AZ QUOTES

Collection Performance



Show me your flowcharts and conceal your tables, and I shall continue to be mystified. Show me your tables, and I won't usually need your flowcharts; they'll be obvious.

— Fred Brooks —

AZ QUOTES

Collection Performance - immutable sequence

	head	tail	apply	update	prepend	append	insert
List	C	C	L	L	C	L	
Stream	C	C	L	L	C	L	
Vector	eC	eC	eC	eC	eC	eC	
Stack	C	C	L	L	C	C	L
Queue	aC	aC	L	L	L	C	
Range	C	C	C				
String	C	L	C	L	L	L	

Collection Performance - mutable sequence

	head	tail	apply	update	prepend	append	insert
ArrayBuffer	C	L	C	C	L	aC	L
ListBuffer	C	L	L	L	C	C	L
StringBuilder	C	L	C	C	L	aC	L
MutableList	C	L	L	L	C	C	L
Queue	C	L	L	L	C	C	L
ArraySeq	C	L	C	C			
Stack	C	L	L	L	C	L	L
ArrayList	C	L	C	C	aC	L	L
Array	C	L	C	C			

Collection Performance - set & map

	lookup	add	remove	min
immutable				
HashSet/HashMap	eC	eC	eC	L
TreeSet/TreeMap	Log	Log	Log	Log
BitSet	C	L	L	eC*
ListMap	L	L	L	L
mutable				
HashSet/HashMap	eC	eC	eC	L
WeakHashMap	eC	eC	eC	L
BitSet	C	aC	C	eC
TreeSet	Log	Log	Log	Log

How to analyze Scala?

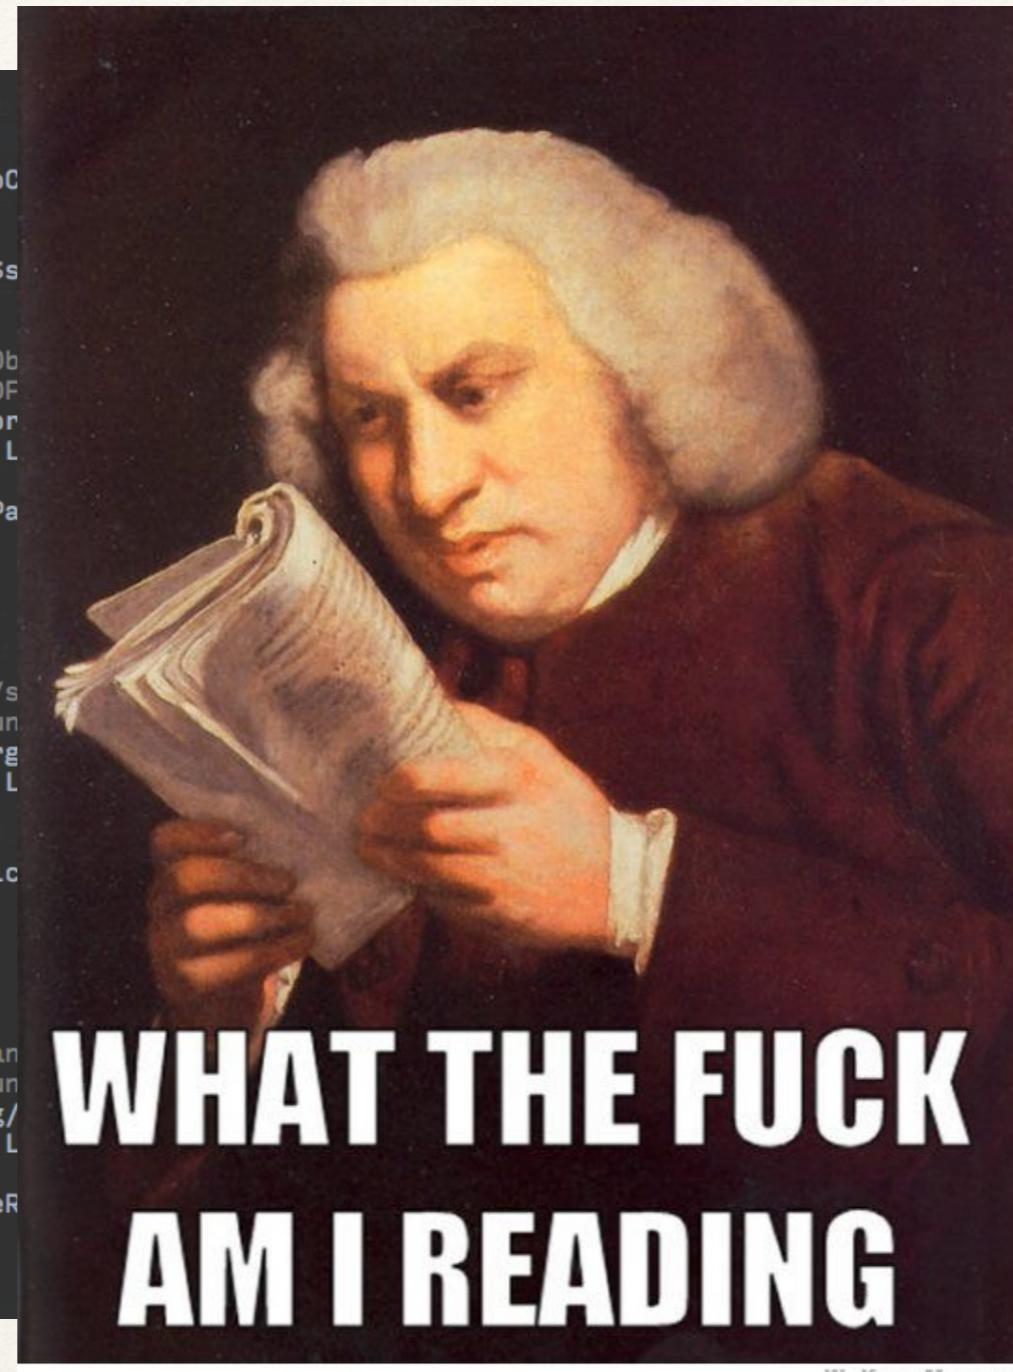
- ❖ Scala will be compiled to bytecode.

Bytecode

```
private Lscala/Option; checkpointData
    if (this.elementClassTag.isAssignableFrom(elementRuntimeClass)) {
        // access flags 0x82
        private transient Z org$apache$spark$rdd$RDD$$doCheckpointCalled
            Lorg/apache/spark/rdd/RDD; doCheckpointCalled(Ljava/lang/Object;)Z
        // access flags 0x82
        private transient Lorg/slf4j/Logger; org$apache$spark$Logging$$log_
            Lorg/apache/spark/Logging; log_(Ljava/lang/Object;Ljava/lang/String;Lscala/runtime/Null$)
        // access flags 0x9
        // signature <K:Ljava/lang/Object;V:Ljava/lang/Object;>(Lorg/apache/spark/rdd/RDD<Lscala/Tuple2<TK;TV;>;>;)Lscala/runtime/Null$;
        // declaration: scala.runtime.Null$ rddToPairRDDFunctions$default$4<K, V>(org.apache.spark.rdd.RDD<scala.Tuple2<K, V>>)
        public static rddToPairRDDFunctions$default$4(Lorg/apache/spark/rdd/RDD;)Lscala/runtime/Null$;
            GETSTATIC org/apache/spark/rdd/RDD$.MODULE$ : Lorg/apache/spark/rdd/RDD$;
            ALOAD 0
            INVOKEVIRTUAL org/apache/spark/rdd/RDD$.rddToPairRDDFunctions$default$4 (Lorg/apache/spark/rdd/RDD;)Lscala/runtime/Null$;
            ARETURN
            MAXSTACK = 2
            MAXLOCALS = 1
        // access flags 0x9
        // signature <T:Ljava/lang/Object;>(Lorg/apache/spark/rdd/RDD<TT;>;Lscala/math/Numeric<TT;>;)Lorg/apache/spark/rdd/DoubleRDDFunctions;
        // declaration: org.apache.spark.rdd.DoubleRDDFunctions numericRDDToDoubleRDDFunctions<T>(org.apache.spark.rdd.RDD<T>, scala.math.Numeric<T>)
        public static numericRDDToDoubleRDDFunctions(Lorg/apache/spark/rdd/RDD;Lscala/math/Numeric;)Lorg/apache/spark/rdd/DoubleRDDFunctions;
            GETSTATIC org/apache/spark/rdd/RDD$.MODULE$ : Lorg/apache/spark/rdd/RDD$;
            ALOAD 0
            ALOAD 1
            INVOKESTATIC Lorg/apache/spark/util/Util;.scalarTachyonFS()Lscala/Function1;
            ALOAD 1
            INVOKEVIRTUAL org/apache/spark/rdd/RDD$.numericRDDToDoubleRDDFunctions (Lorg/apache/spark/rdd/RDD;Lscala/math/Numeric;)Lorg/apache/spark/rdd/DoubleRDDFunctions;
            ARETURN
            MAXSTACK = 3
            MAXLOCALS = 2
        // access flags 0x9
        // signature (Lorg/apache/spark/rdd/RDD<Ljava/lang/Object;>;)Lorg/apache/spark/rdd/DoubleRDDFunctions;
        // declaration: org.apache.spark.rdd.DoubleRDDFunctions doubleRDDToDoubleRDDFunctions(org.apache.spark.rdd.RDD<java.lang.Object>)
        public static doubleRDDToDoubleRDDFunctions(Lorg/apache/spark/rdd/RDD;)Lorg/apache/spark/rdd/DoubleRDDFunctions;
            GETSTATIC org/apache/spark/rdd/RDD$.MODULE$ : Lorg/apache/spark/rdd/RDD$;
            ALOAD 0
            INVOKEVIRTUAL org/apache/spark/rdd/RDD$.doubleRDDToDoubleRDDFunctions (Lorg/apache/spark/rdd/RDD;)Lorg/apache/spark/rdd/DoubleRDDFunctions;
            ARETURN
            MAXSTACK = 2
            MAXLOCALS = 1
```

Bytecode

```
private Lscala/Option; checkpointData  
// access flags 0x82  
private transient Z org$apache$spark$rdd$RDD$$doC  
  
// access flags 0x82  
private transient Lorg.slf4j/Logger; org$apache$s  
  
// access flags 0x9  
// signature <K:Ljava/lang/Object;V:Ljava/lang/Ob  
// declaration: scala.runtime.Null$ rddToPairRDDF  
public static rddToPairRDDFunctions$default$4(Lor  
    GETSTATIC org/apache/spark/rdd/RDD$.MODULE$ : L  
    ALOAD 0  
    INVOKEVIRTUAL org/apache/spark/rdd/RDD$.rddToPa  
    ARETURN  
    MAXSTACK = 2  
    MAXLOCALS = 1  
  
// access flags 0x9  
// signature <T:Ljava/lang/Object;>(Lorg/apache/s  
// declaration: org.apache.spark.DoubleRDDFun  
public static numericRDDToDoubleRDDFunctions(Lor  
    GETSTATIC org/apache/spark/rdd/RDD$.MODULE$ : L  
    ALOAD 0  
    ALOAD 1  
    INVOKEVIRTUAL org/apache/spark/rdd/RDD$.numeric  
    ARETURN  
    MAXSTACK = 3  
    MAXLOCALS = 2  
  
// access flags 0x9  
// signature (Lorg/apache/spark/rdd/RDD<Ljava/lan  
// declaration: org.apache.spark.DoubleRDDFun  
public static doubleRDDToDoubleRDDFunctions(Lor  
    GETSTATIC org/apache/spark/rdd/RDD$.MODULE$ : L  
    ALOAD 0  
    INVOKEVIRTUAL org/apache/spark/rdd/RDD$.doubleR  
    ARETURN  
    MAXSTACK = 2  
    MAXLOCALS = 1
```



```
a/runtime/Null$;  
K, V>>)  
  
time/Null$;  
  
dd/DoubleRDDFunctions;  
DD<T>, scala.math.Numeric<T>)  
d/DoubleRDDFunctions;  
  
information.  
Numeric;)Lorg/apache/spark/rdd/DoubleRDDFunctions;  
  
ava.lang.Object>)  
  
spark/rdd/DoubleRDDFunctions;
```

Reverse Engineering

- ❖ scalac -Xprint:all {source code}
- ❖ Java Decompiler
- ❖ Luyten
- ❖ But Java bytecode is still the best way to know under the hood.

combinator vs recursion vs iteration

- ❖ There is no free lunch!

How to benchmark?

- ❖ sbt-jmh
- ❖ Benchmark Environment
 - ❖ Debian 8
 - ❖ OpenJDK 7
 - ❖ Scala 2.10.5

Loop2

[info] Benchmark	Mode	Cnt	Score	Error	Units
[info] Loop2.combinator	thrpt	400	58.372 ± 0.173	0.173	ops/ms
[info] Loop2.forLoop	thrpt	400	77.289 ± 0.992	0.992	ops/ms
[info] Loop2.tailCall	thrpt	400	168.048 ± 5.360	5.360	ops/ms
[info] Loop2.whileLoop	thrpt	400	202.866 ± 4.149	4.149	ops/ms
[info] Loop2.combinator	avgt	400	0.017 ± 0.001	0.001	ms/op
[info] Loop2.forLoop	avgt	400	0.013 ± 0.001	0.001	ms/op
[info] Loop2.tailCall	avgt	400	0.006 ± 0.001	0.001	ms/op
[info] Loop2.whileLoop	avgt	400	0.006 ± 0.001	0.001	ms/op
[info] Loop2.combinator	ss	400	0.046 ± 0.009	0.009	ms/op
[info] Loop2.forLoop	ss	400	0.038 ± 0.004	0.004	ms/op
[info] Loop2.tailCall	ss	400	0.067 ± 0.007	0.007	ms/op
[info] Loop2.whileLoop	ss	400	0.067 ± 0.007	0.007	ms/op

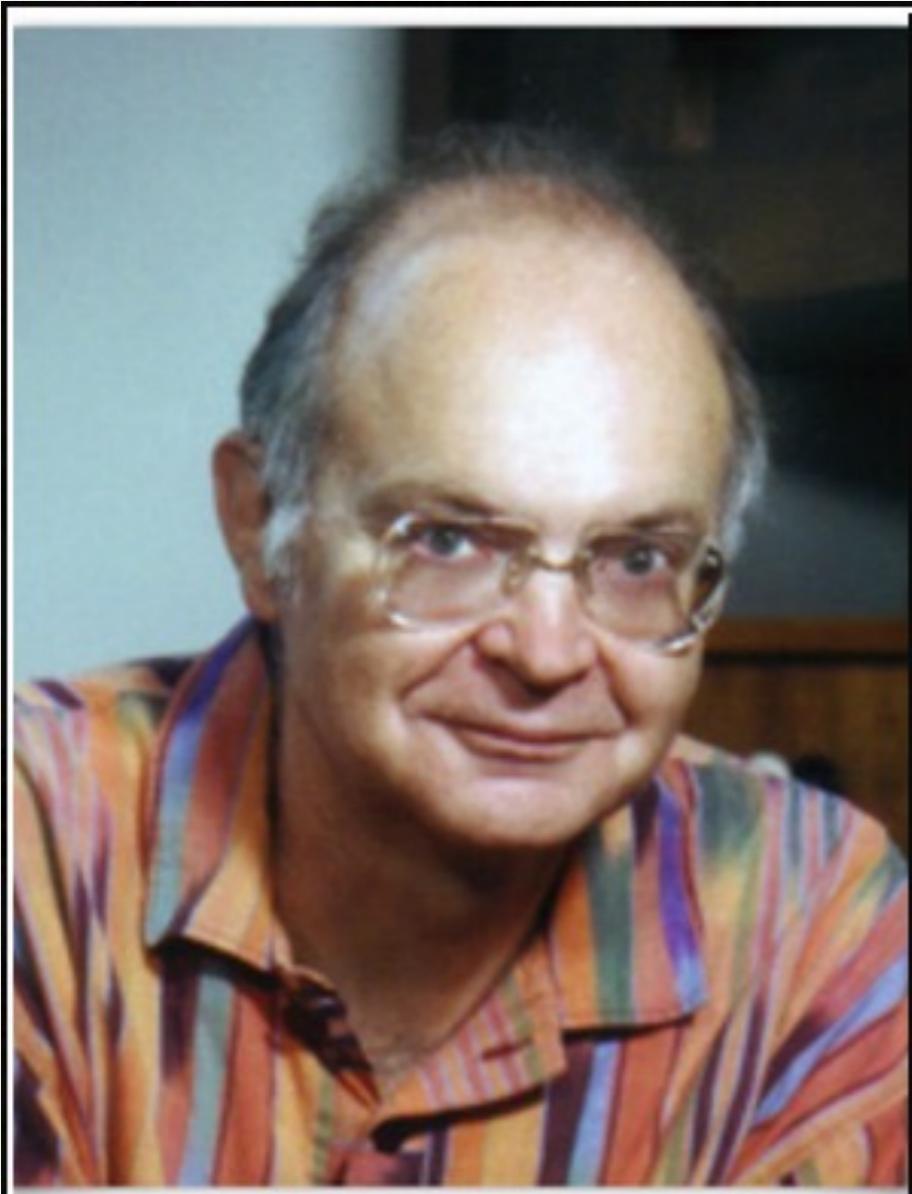
Loop3

[info] Benchmark	Mode	Cnt	Score	Error	Units
[info] Loop3.combinator	thrpt	400	57.734 ± 0.216	0.216	ops/ms
[info] Loop3.forLoopListBuffer	thrpt	400	70.846 ± 0.363	0.363	ops/ms
[info] Loop3.tailCallIfElse	thrpt	400	73.723 ± 0.190	0.190	ops/ms
[info] Loop3.tailCallPatternMatch	thrpt	400	79.851 ± 0.393	0.393	ops/ms
[info] Loop3.combinator	avgt	400	0.017 ± 0.001	0.001	ms/op
[info] Loop3.forLoopListBuffer	avgt	400	0.014 ± 0.001	0.001	ms/op
[info] Loop3.tailCallIfElse	avgt	400	0.014 ± 0.001	0.001	ms/op
[info] Loop3.tailCallPatternMatch	avgt	400	0.013 ± 0.001	0.001	ms/op
[info] Loop3.combinator	ss	400	0.039 ± 0.001	0.001	ms/op
[info] Loop3.forLoopListBuffer	ss	400	0.030 ± 0.001	0.001	ms/op
[info] Loop3.tailCallIfElse	ss	400	0.033 ± 0.001	0.001	ms/op
[info] Loop3.tailCallPatternMatch	ss	400	0.029 ± 0.001	0.001	ms/op

Loop4

[info] Benchmark	Mode	Cnt	Score	Error	Units
[info] Loop4.combinator	thrpt	400	5.026 ± 0.019	0.019	ops/ms
[info] Loop4.tailCall	thrpt	400	86.751 ± 0.400	0.400	ops/ms
[info] Loop4.whileLoop	thrpt	400	77.319 ± 2.406	2.406	ops/ms
[info] Loop4.combinator	avgt	400	0.199 ± 0.001	0.001	ms/op
[info] Loop4.tailCall	avgt	400	0.012 ± 0.001	0.001	ms/op
[info] Loop4.whileLoop	avgt	400	0.014 ± 0.001	0.001	ms/op
[info] Loop4.combinator	ss	400	1.154 ± 0.106	0.106	ms/op
[info] Loop4.tailCall	ss	400	0.037 ± 0.007	0.007	ms/op
[info] Loop4.whileLoop	ss	400	0.071 ± 0.014	0.014	ms/op

combinator vs recursion vs loop

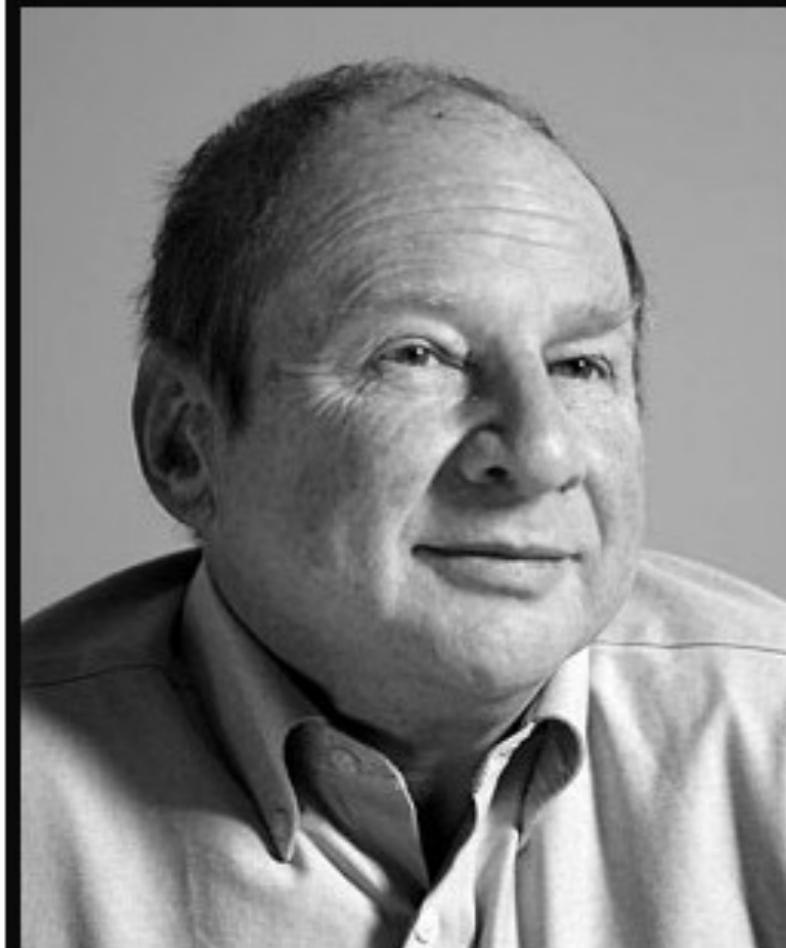


Premature optimization is the root
of all evil.

— Donald Knuth —

AZ QUOTES

combinator vs recursion vs loop



Programs must be written for people to read,
and only incidentally for machines to execute.

(Hal Abelson)

izquotes.com

combinator vs recursion vs loop

- ❖ Consider using combinators first.
- ❖ If this becomes too tedious, or efficiency is a big concern, fall back on tail-recursive functions.
- ❖ Loop can be used in simple case, or when the computation inherently modifies state.
- ❖ from Martin Odersky: Scala with Style

loop vs recursion vs combinators

- ❖ Summary:
 - ❖ Choice the way has **better readability** and it's easier to reason.
 - ❖ Based on my summary, it seems my talk is useless.
 - ❖ No.
 - ❖ Choice hot spot to optimise!
 - ❖ Yes.
 - ❖ The most important thing is **Scala compiler/JVM will evolve**.

Loop5

[info] Benchmark	Mode	Cnt	Score	Error	Units
[info] Loop5.forLoopDouble	thrpt	400	$\approx 10^{-4}$		ops/us
[info] Loop5.forLoopFloat	thrpt	400	$\approx 10^{-4}$		ops/us
[info] Loop5.forLoopInt	thrpt	400	$\approx 10^{-4}$		ops/us
[info] Loop5.forLoopLong	thrpt	400	$\approx 10^{-4}$		ops/us
[info] Loop5.whileLoopDouble	thrpt	400	$\approx 10^{-4}$		ops/us
[info] Loop5.whileLoopFloat	thrpt	400	$\approx 10^{-4}$		ops/us
[info] Loop5.whileLoopInt	thrpt	400	$\approx 10^{-4}$		ops/us
[info] Loop5.whileLoopLong	thrpt	400	$\approx 10^{-4}$		ops/us
[info] Loop5.forLoopDouble	avgt	400	20889.321 \pm 54.857		us/op
[info] Loop5.forLoopFloat	avgt	400	12479.768 \pm 295.988		us/op
[info] Loop5.forLoopInt	avgt	400	9532.830 \pm 17.743		us/op
[info] Loop5.forLoopLong	avgt	400	20062.058 \pm 13.270		us/op
[info] Loop5.whileLoopDouble	avgt	400	20222.500 \pm 102.163		us/op
[info] Loop5.whileLoopFloat	avgt	400	11553.178 \pm 274.842		us/op
[info] Loop5.whileLoopInt	avgt	400	9488.849 \pm 7.718		us/op
[info] Loop5.whileLoopLong	avgt	400	19047.508 \pm 4.433		us/op
[info] Loop5.forLoopDouble	ss	400	20959.765 \pm 46.540		us/op
[info] Loop5.forLoopFloat	ss	400	11780.323 \pm 106.223		us/op
[info] Loop5.forLoopInt	ss	400	9586.689 \pm 31.832		us/op
[info] Loop5.forLoopLong	ss	400	19077.921 \pm 28.002		us/op
[info] Loop5.whileLoopDouble	ss	400	20046.600 \pm 5.705		us/op
[info] Loop5.whileLoopFloat	ss	400	10731.143 \pm 9.356		us/op
[info] Loop5.whileLoopInt	ss	400	9513.217 \pm 8.281		us/op
[info] Loop5.whileLoopLong	ss	400	19078.069 \pm 12.329		us/op

Loop5 - Scala vs Java

	Scala(us/op)	Java(us/op)
for-loop-int	9532.83	9489.191
while-loop-int	9488.849	9491.391
for-loop-long	20062.058	18991.304
while-loop-long	19047.508	18988.561
for-loop-float	12479.768	11365.767
while-loop-float	11553.178	11198.945
for-loop-double	20889.321	19798.314
while-loop-double	20222.5	19772.461

Loop5 - Scala vs Java

```
def forLoopInt(state:  
BenchmarkState) = {  
    var s = 0  
  
    for (i <- 1 to state.n;  
         j <- 1 to state.n;  
         k <- 1 to state.n)  
        s += state.random.nextInt(4)  
    s  
}
```

```
public int forLoopInt(BenchmarkStateLoop5Java  
state) {  
    int s = 0;  
  
    for (int i = 1; i <= state.n; i++)  
        for (int j = 1; j <= state.n; j++)  
            for (int k = 1; k <= state.n; k++)  
                s += state.random.nextInt(4);  
  
    return s;  
}
```

Loop5 - Scala vs Java

- ❖ Scala nested for-loop is just slightly slower than Scala while-loop.
- ❖ Scala while-loop is pretty close to Java for-loop and Java while-loop.

primitive type vs boxed primitive

[info] Specialization.tuple2LoopDouble	thrpt	400	0.006 ± 0.001	ops/ms
[info] Specialization.tuple2LoopInt	thrpt	400	0.006 ± 0.001	ops/ms
[info] Specialization.tuple2SpeicalLoopDouble	thrpt	400	0.007 ± 0.001	ops/ms
[info] Specialization.tuple2SpeicalLoopInt	thrpt	400	0.008 ± 0.001	ops/ms
[info] Specialization.tuple2LoopDouble	avgt	400	152.308 ± 1.626	ms/op
[info] Specialization.tuple2LoopInt	avgt	400	154.144 ± 1.676	ms/op
[info] Specialization.tuple2SpeicalLoopDouble	avgt	400	155.701 ± 2.723	ms/op
[info] Specialization.tuple2SpeicalLoopInt	avgt	400	133.234 ± 1.106	ms/op
[info] Specialization.tuple2LoopDouble	ss	400	155.689 ± 2.899	ms/op
[info] Specialization.tuple2LoopInt	ss	400	155.466 ± 1.683	ms/op
[info] Specialization.tuple2SpeicalLoopDouble	ss	400	156.365 ± 3.034	ms/op
[info] Specialization.tuple2SpeicalLoopInt	ss	400	132.948 ± 1.291	ms/op

primitive type vs boxed primitive

[info] Specialization.tuple2LoopDouble	thrpt	400	0.006 ± 0.001	ops/ms
[info] Specialization.tuple2LoopInt	thrpt	400	0.006 ± 0.001	ops/ms
[info] Specialization.tuple2SpeicalLoopDouble	thrpt	400	0.007 ± 0.001	ops/ms
[info] Specialization.tuple2SpeicalLoopInt	thrpt	400	0.008 ± 0.001	ops/ms
[info] Specialization.tuple2LoopDouble	avgt	400	152.308 ± 1.626	ms/op
[info] Specialization.tuple2LoopInt	avgt	400	154.144 ± 1.676	ms/op
[info] Specialization.tuple2SpeicalLoopDouble	avgt	400	155.701 ± 2.723	ms/op
[info] Specialization.tuple2SpeicalLoopInt	avgt	400	133.234 ± 1.106	ms/op
[info] Specialization.tuple2LoopDouble	ss	400	155.689 ± 2.899	ms/op
[info] Specialization.tuple2LoopInt	ss	400	155.466 ± 1.683	ms/op
[info] Specialization.tuple2SpeicalLoopDouble	ss	400	156.365 ± 3.034	ms/op
[info] Specialization.tuple2SpeicalLoopInt	ss	400	132.948 ± 1.291	ms/op

- ❖ I didn't add annotation for Double, so it still does lots of box and unbox.

primitive type vs boxed primitive

- ❖ Primitive type always faster than boxed primitive.
- ❖ Multiplying two matrices of 200x100 is **2.5 times** faster with specialization.
- ❖ from Specialization in Scala 2.8

String Size

- ❖ Object header is around 16 bytes in 64-bit and 12 bytes in 32-bit.
- ❖ String has 40 bytes (32-bit) or 56 bytes (64-bit) overhead.
 - ❖ Apache spark suggests that uses numeric IDs or enumeration other than String.
- ❖ Reduce memory overhead.

Why Scala?

- ❖ Why FP?
- ❖ Benchmark
 - ❖ Speed
 - ❖ Memory footprint
 - ❖ Line of code
 - ❖ Size of Binary and JAR
 - ❖ Compilation time
 - ❖ Go vs Scala in parallel and concurrent computing
- ❖ Subjective opinion

Why FP(Scala)?

- ❖ Fewer errors
- ❖ Better modularity
- ❖ High-level abstractions
- ❖ Shorter code
- ❖ Increased developer productivity(*)
- ❖ Parallel (or Concurrent) computing(*)
- ❖ Distributed computing(*)

Why Scala? (Speed)

Benchmark	Time [sec]	Factor
C++ Opt	23	1.0x
C++ Dbg	197	8.6x
Java 64-bit	134	5.8x
Java 32-bit	290	12.6x
Java 32-bit GC*	106	4.6x
Java 32-bit SPEC GC	89	3.7x
Scala	82	3.6x
Scala low-level*	67	2.9x
Scala low-level GC*	58	2.5x
Go 6g	161	7.0x
Go Pro*	126	5.5x

Why Scala? (Speed)

- ❖ We find that in regards to performance, C++ wins out by a large margin. However, it also required **the most extensive tuning efforts**, many of which were done at a level of sophistication that would **not be available to the average programmer**.
- ❖ from [Loop Recognition in C++/Java/Go/Scala by Google](#)

Why Scala? (Memory)

Benchmark	Virt	Real	Factor Virt	Factor Real
C++ Opt	184m	163m	1.0	1.0
C++ Dbg	474m	452m	2.6-3.0	2.8
Java	1109m	617m	6.0	3.7
Scala	1111m	293m	6.0	1.8
Go	16.2g	501m	90	3.1

Why Scala? (Lines of code)

Benchmark	wc -1	Factor
C++ Dbg/Opt	850	1.3x
Java	1068	1.6x
Java Pro	1240	1.9x
Scala	658	1.0x
Scala Pro	297	0.5x
Go	902	1.4x
Go Pro	786	1.2x

Why Scala?

- ❖ Scala concise notation and powerful language features allowed for the best optimization of code complexity.
- ❖ from Loop Recognition in C++/Java/Go/Scala by Google

Why Scala? (Parallel)

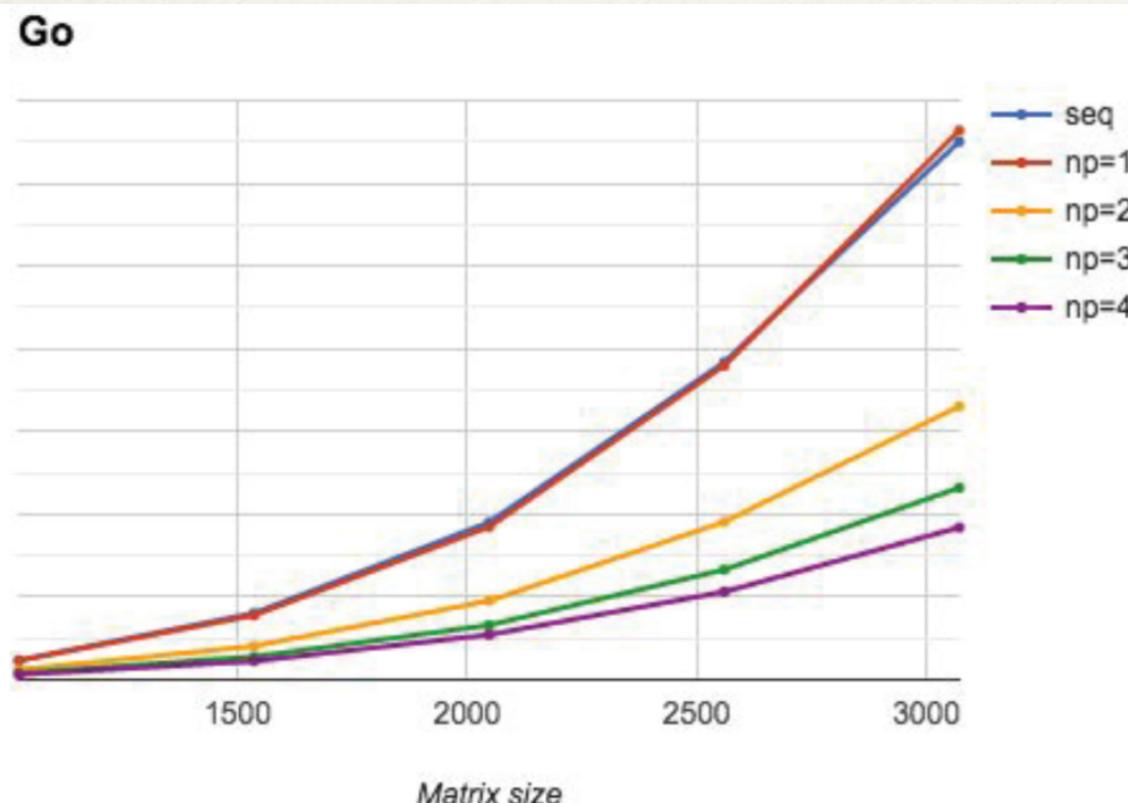


Figure 5-1: Matrix multiplication in Go

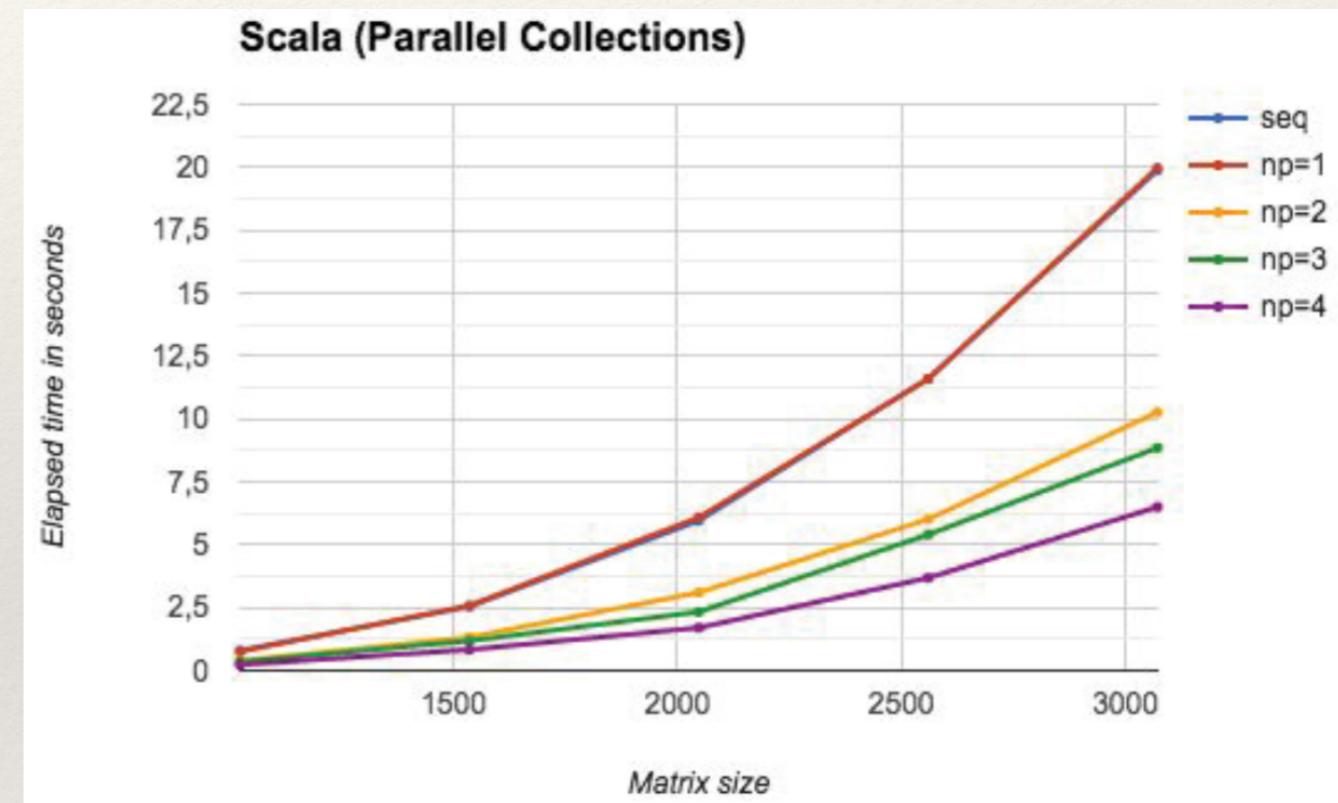


Figure 5-2: Matrix multiplication in Scala with parallel collections

Why Scala? (Concurrent)

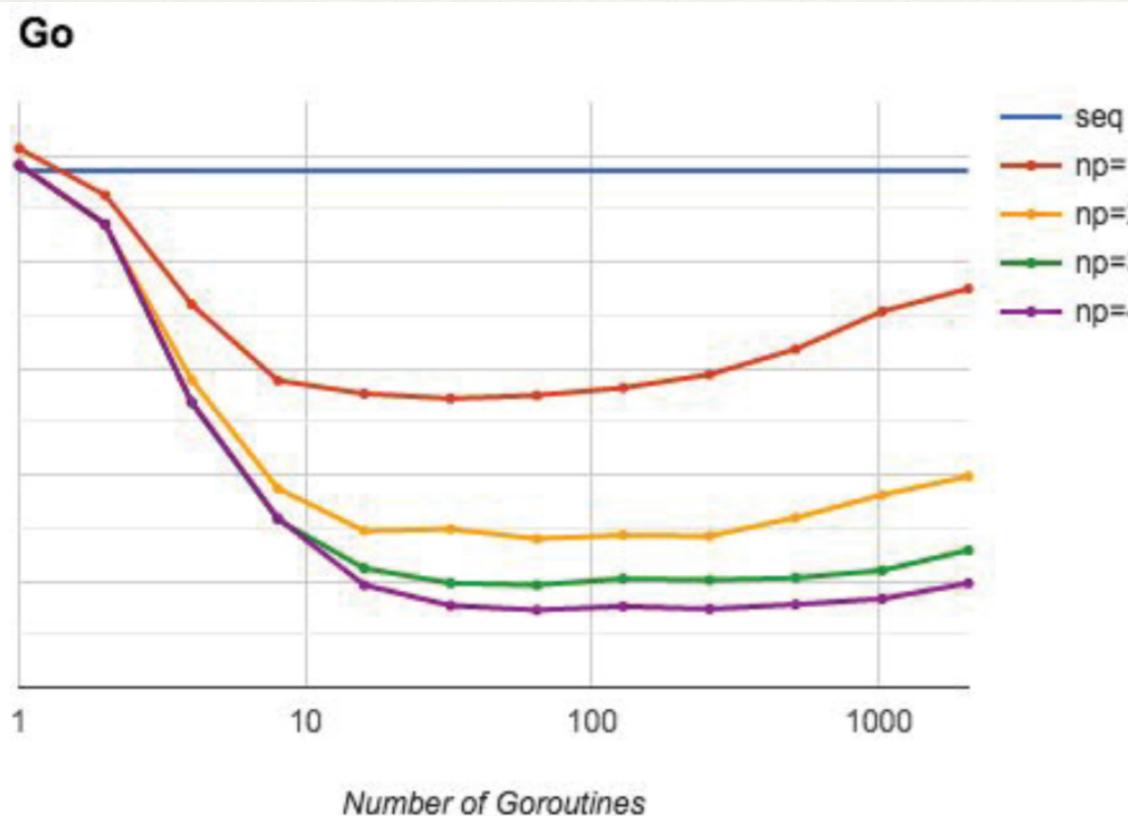


Figure 5-7: Matrix chain multiplication in Go

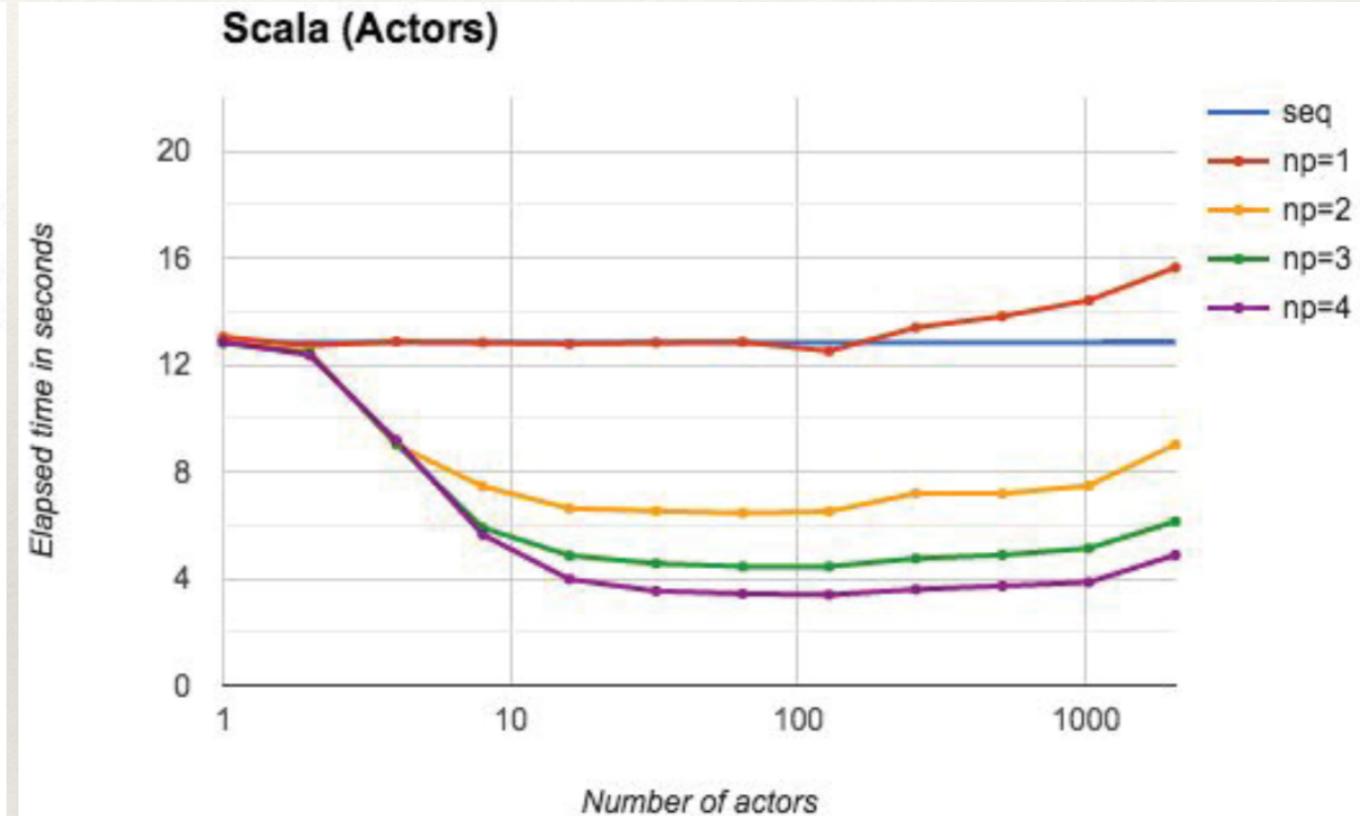


Figure 5-8: Matrix chain multiplication with actors in Scala

Downside of Scala

Benchmark	Binary or Jar [Byte]	Factor
C++ Dbg	592892	45x
C++ Opt	41507	3.1x
Java	13215	1.0x
Java Pro	21047	1.6x
Scala	48183	3.6x
Scala Pro	36863	2.8x
Go	1249101	94x
Go Pro	1212100	92x

Downside of Scala

Benchmark	Compile Time	Factor
C++ Dbg	3.9	6.5x
C++ Opt	3.0	5.0x
Java	3.1	5.2x
Java Pro	3.0	5.0x
Scala scalac	13.9	23.1x
Scala fsc	3.8	6.3x
Scala Pro scalac	11.3	18.8x
Scala Pro fsc	3.5	5.8x
Go	1.2	2.0x
Go Pro	0.6	1.0x

Why Scala? (Subjective)



A language that doesn't affect the way you think about programming is not worth knowing.

— *Alan Perlis* —

AZ QUOTES

Downside of Scala (Subjective)

- ❖ I enjoy writing Scala code but I hate reading other people's Scala code. from Reddit
 - ❖ Code review could solve this issue.
- ❖ OOP or FP?
 - ❖ double-edged sword

Why Scala? (Summary)

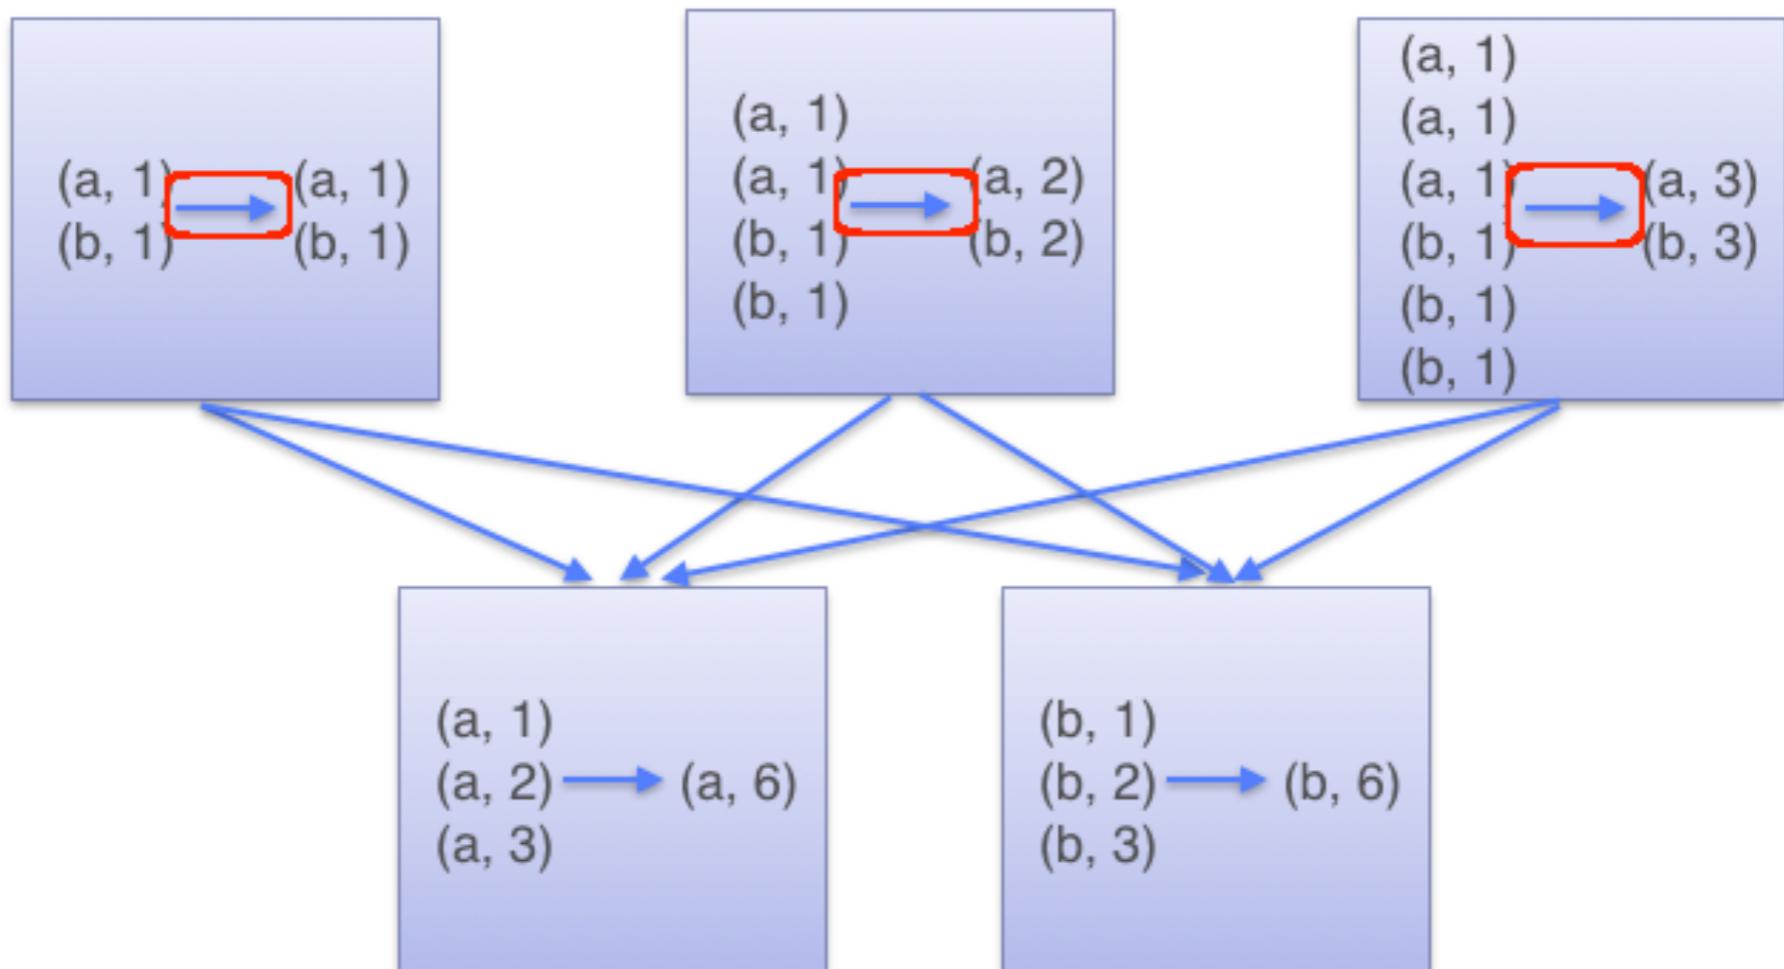
- ❖ Scala **isn't the best** at any of these things though. It's great because it's an **all-rounder** language.
- ❖ Scala doesn't - quite - have the safety and conciseness of *Haskell*. But it's close.
- ❖ Scala doesn't - quite - have the enterprise support and tooling infrastructure (monitoring/instrumentation, profiling, debugging, IDEs, library ecosystem) that *Java* does. But it's close.
- ❖ Scala doesn't - quite - have the clarity of *Python*. But it's close.
- ❖ Scala doesn't - quite - have the performance of *C++*. But it's close.
- ❖ From <https://news.ycombinator.com/item?id=9398911> by Michael Donaghy(m50d)

Spark Performance

- ❖ 1. Programming API
 - ❖ reduceByKey
 - ❖ aggregateByKey
 - ❖ coalesce
- ❖ 2. Spark configuration and JVM tuning.
 - ❖ Spark: Java / Kyro, Memory Management, Persist
 - ❖ VM Options: JIT compiler, Memory, GC, JMX, Misc

reduceByKey

ReduceByKey



aggregateByKey

- ❖ Spark Source Code.
 - ❖ Serialise the zero value to a byte array so that we can get a new **clone** of it on each key.
 - ❖ It's a shallow copy so it's faster than new object each time.

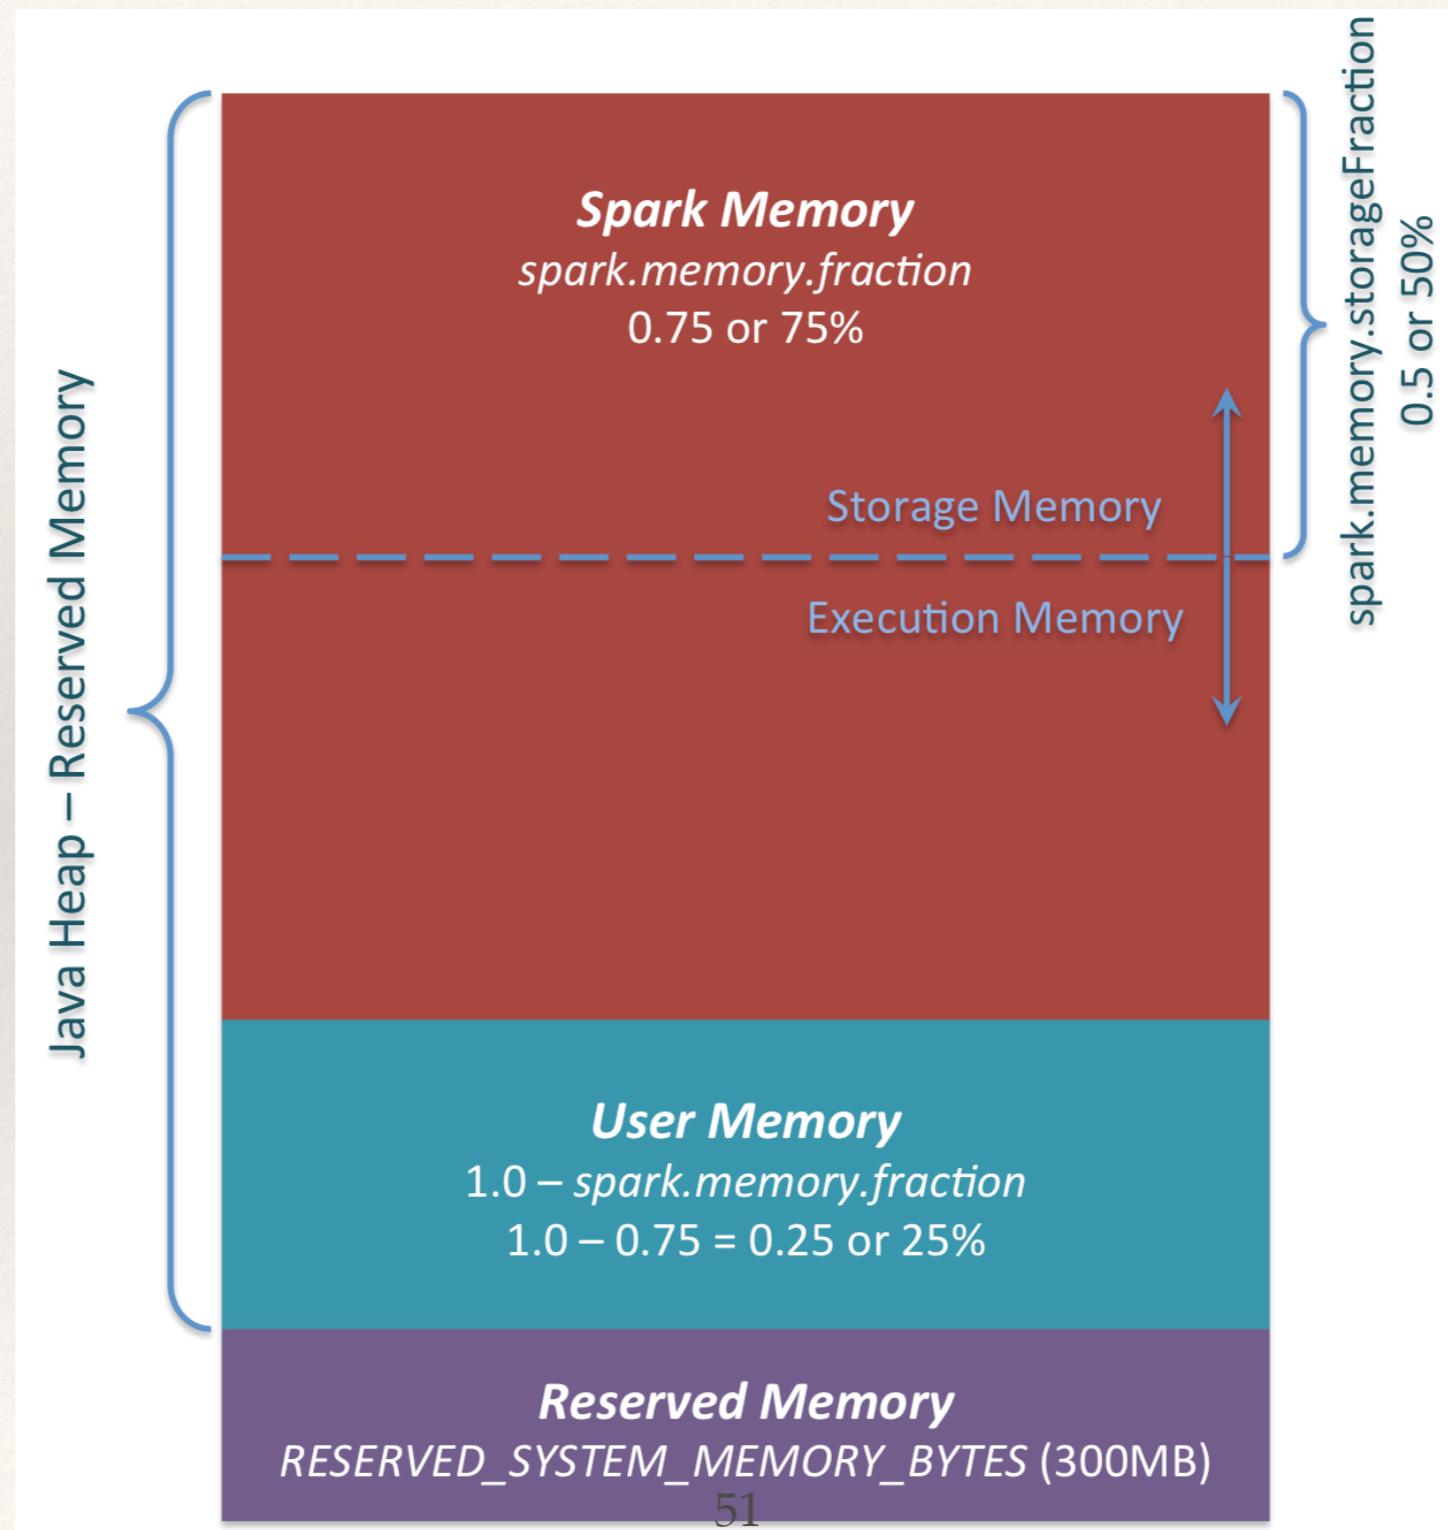
coalesce

- ❖ coalesce() that allows avoiding data movement, but only if you are **decreasing** the number of RDD partitions.
from Learning Spark - Lightning-Fast Big Data Analysis
- ❖ coalesce with shuffle?
 - ❖ If you have a few partitions being **abnormally large**.

Spark Serialization

- ❖ Java serialization
- ❖ Kyro serialization
 - ❖ Pro: Faster than Java serialization .
 - ❖ Con: Not support all Serializable type.

Spark Memory Management



Spark Persist

- ❖ MEMORY_ONLY = Cache
- ❖ MEMORY_SER: Cache with serialization.
- ❖ MEMORY_AND_DISK: Cache, but it will flush to disk when meet the limit.
- ❖ MEMORY_AND_DISK_SER
- ❖ DISK_ONLY
- ❖ OFF_HEAP: Work with Alluxio(Tachyon)
- ❖ *_2: With replication

Spark Persist

- ❖ MEMORY_ONLY: If you have sufficient memory.
- ❖ MEMORY_ONLY_SER: It can help to reduce memory usage. But it also consumes cpu. (Non-serialization will 2~3 time bigger than serialization)
- ❖ *_DISK: Don't use it unless **your computing are expensive.** (Official). I think It will depend on which **type of disk** you use?
- ❖ Remember **unpersist** cache when you don't need.
 - ❖ blocking or not? Blocking for preventing memory pressure.
 - ❖ non-blocking only for you have sufficient memory

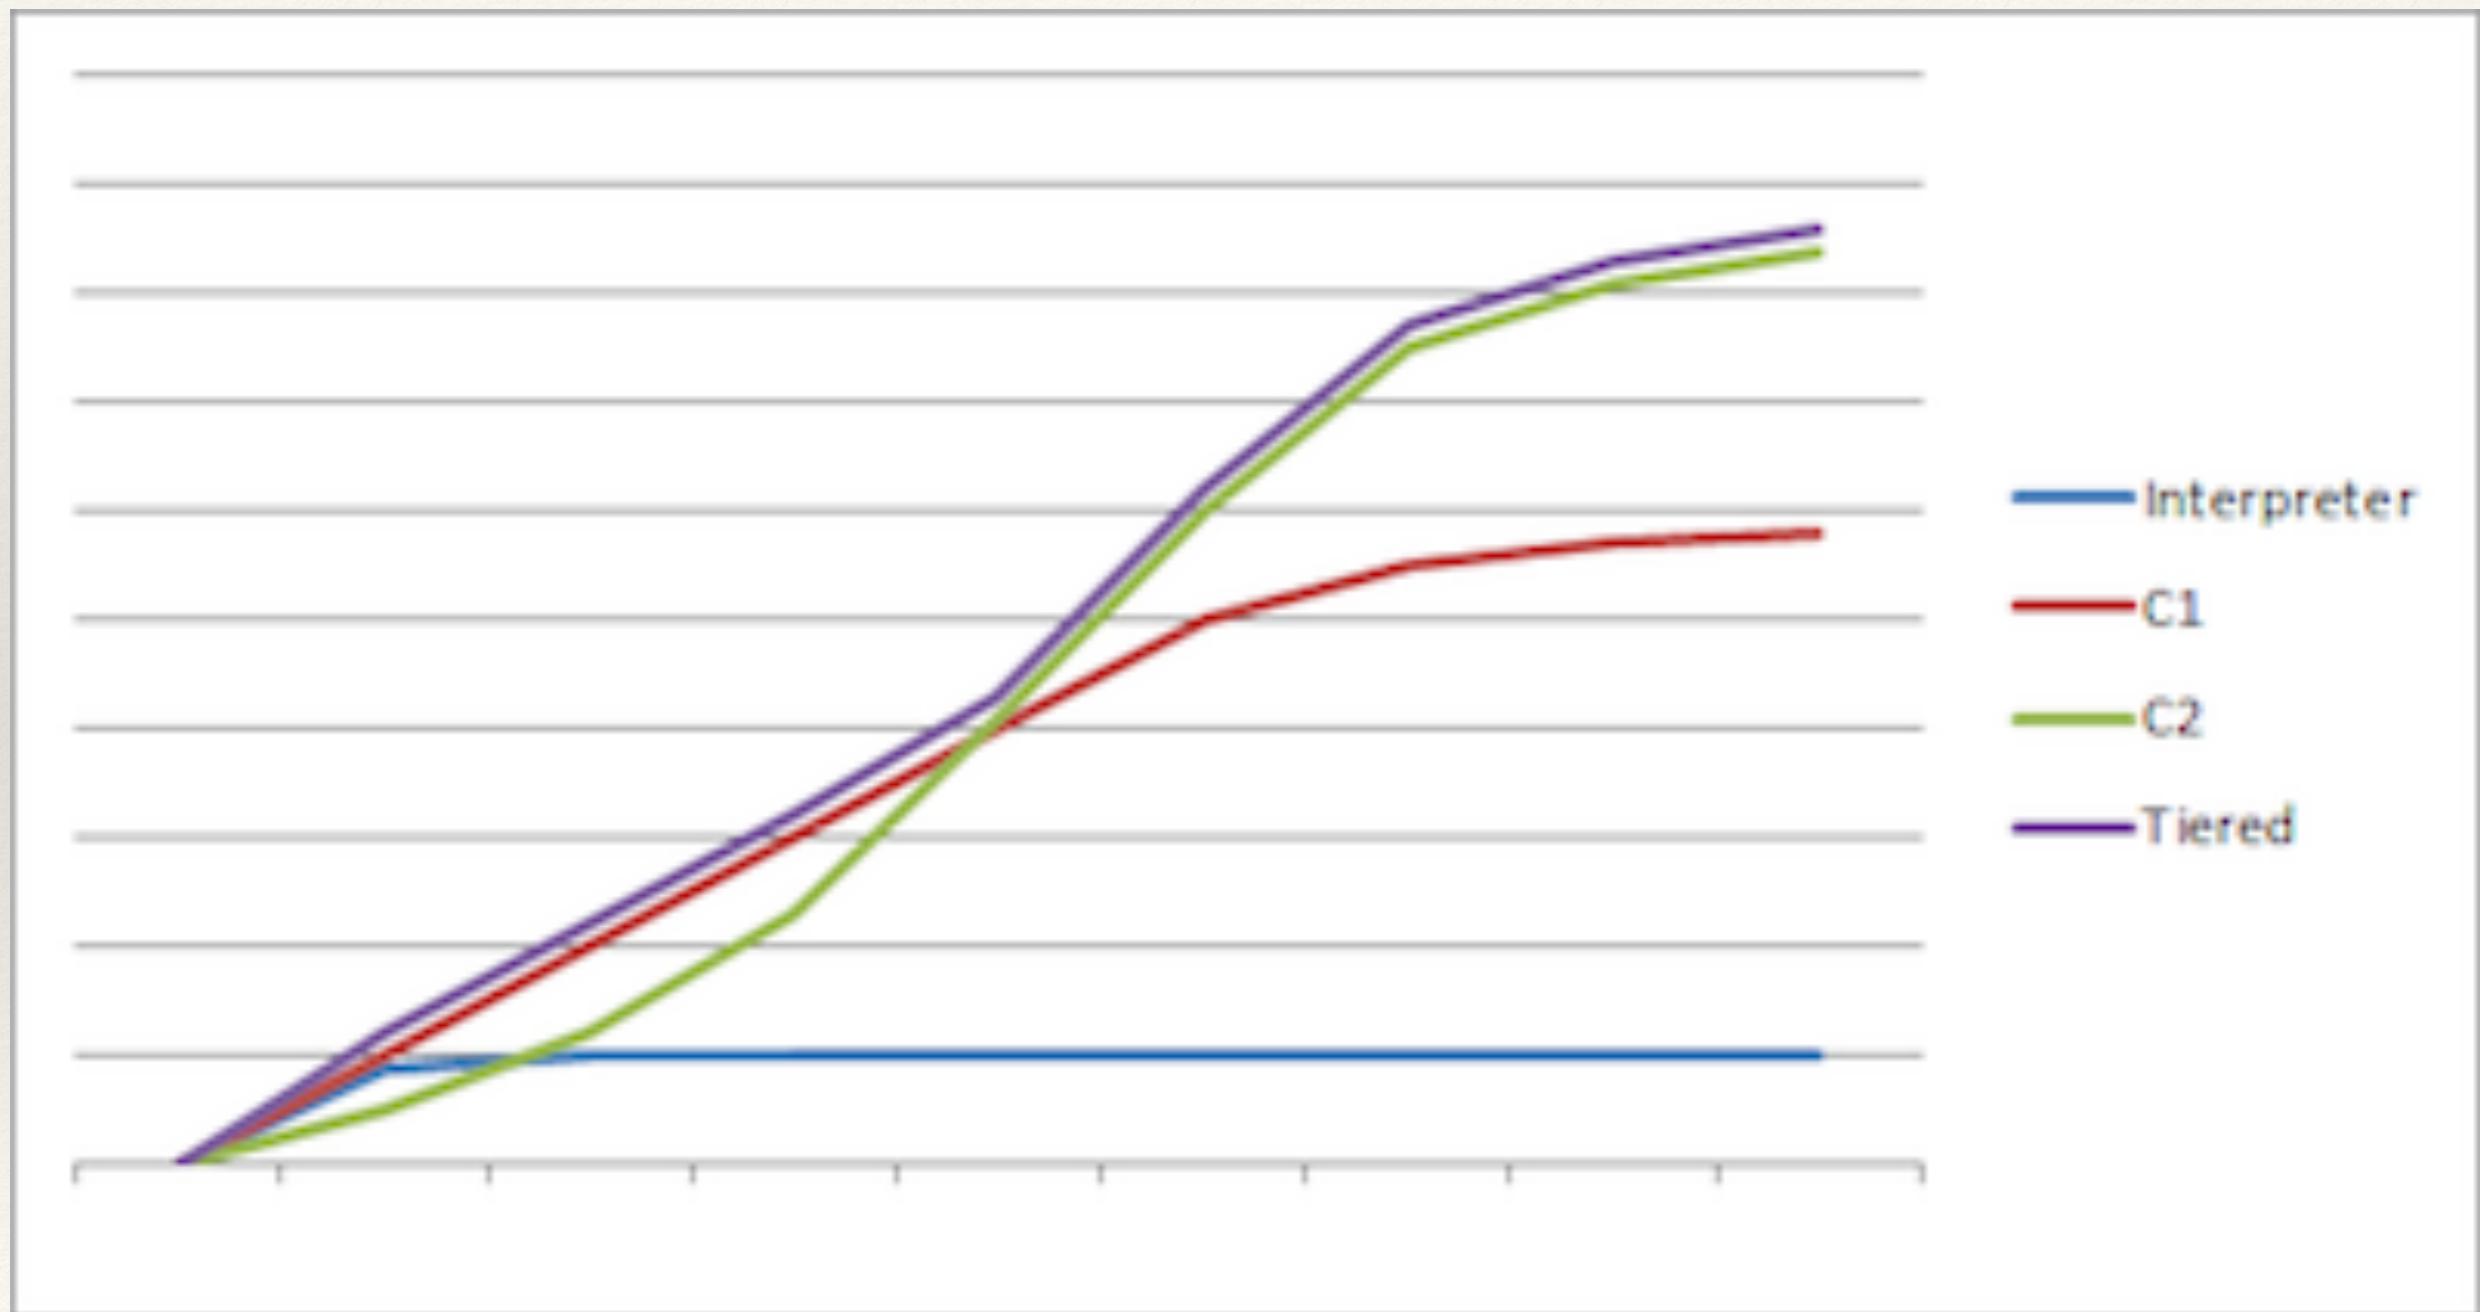
Spark Memory Management

- ❖ How much memory you should allocate?
 - ❖ When JVM uses more than **200G**, you should consider to use 2 workers per node then each work has 100G.
 - ❖ Total Memory * **$3/4$** for Spark. The rest is used for OS and buffer cache.

VM Options - JIT

- ❖ JIT Compiler
 - ❖ -client (1500)
 - ❖ -server (10,000)
 - ❖ -XX:+TieredCompilation
 - ❖ -client (2000) + -server (15,000) from Azul Systems
 - ❖ default setup in JAVA 8.

VM Options - JIT



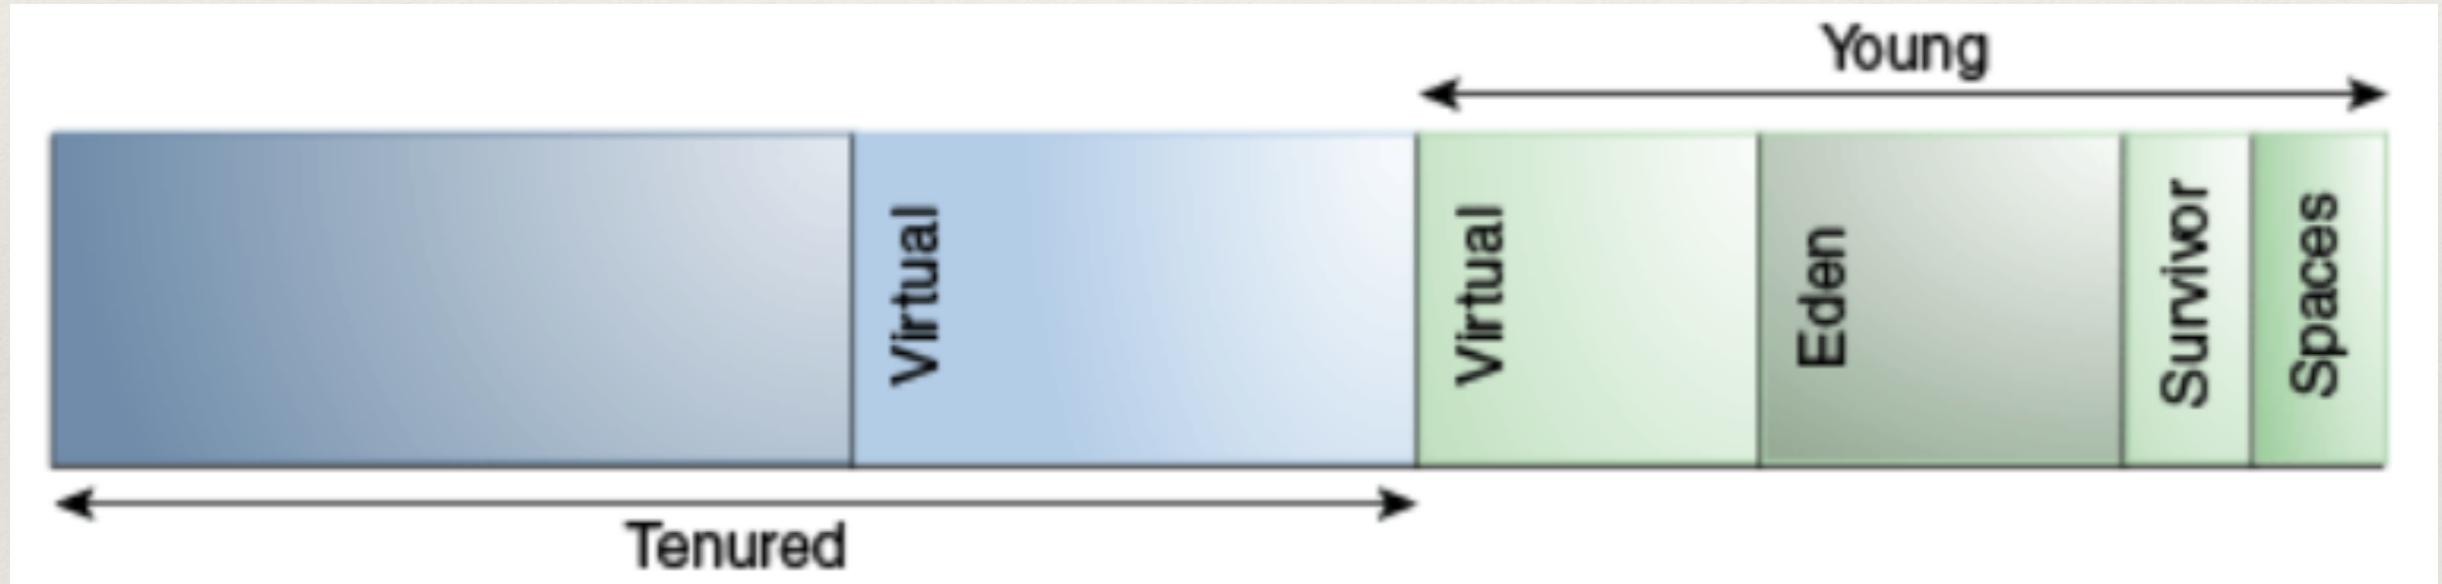
VM Options - Memory

- ❖ `-XX:+UseCompressedOops`
 - ❖ oop = ordinary object pointer
 - ❖ Emulates 35-bit pointers for object references.
 - ❖ When you use memory lower than 32 G.
 - ❖ Reduce memory footprint slightly.

VM Options - GC

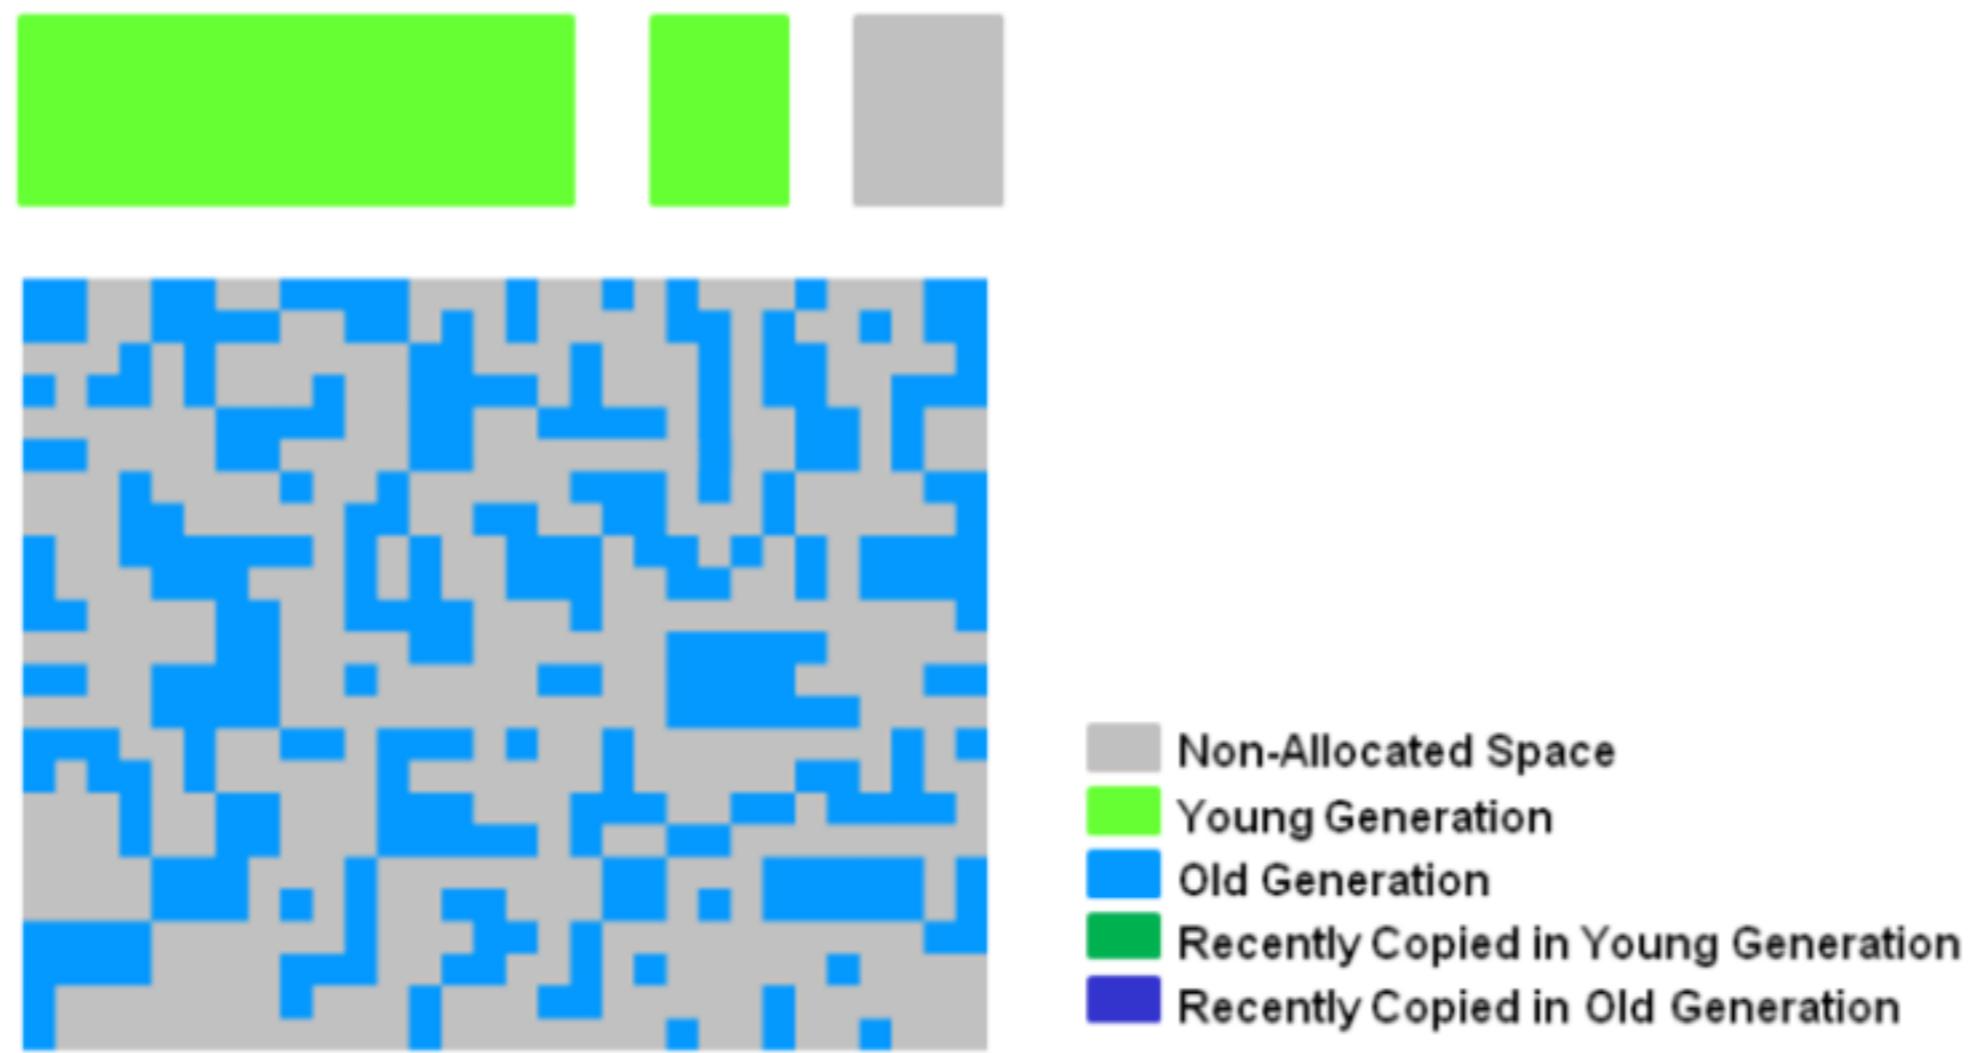
- ❖ Parallel GC
 - ❖ Throughput
 - ❖ Batch, Scientific computing
- ❖ CMS GC
 - ❖ Latency
 - ❖ Streaming, Web Service
- ❖ G1 GC
 - ❖ Throughput & Latency
 - ❖ Batch & Streaming
 - ❖ The long term replacement for CMS.

VM Options - Parallel Collector



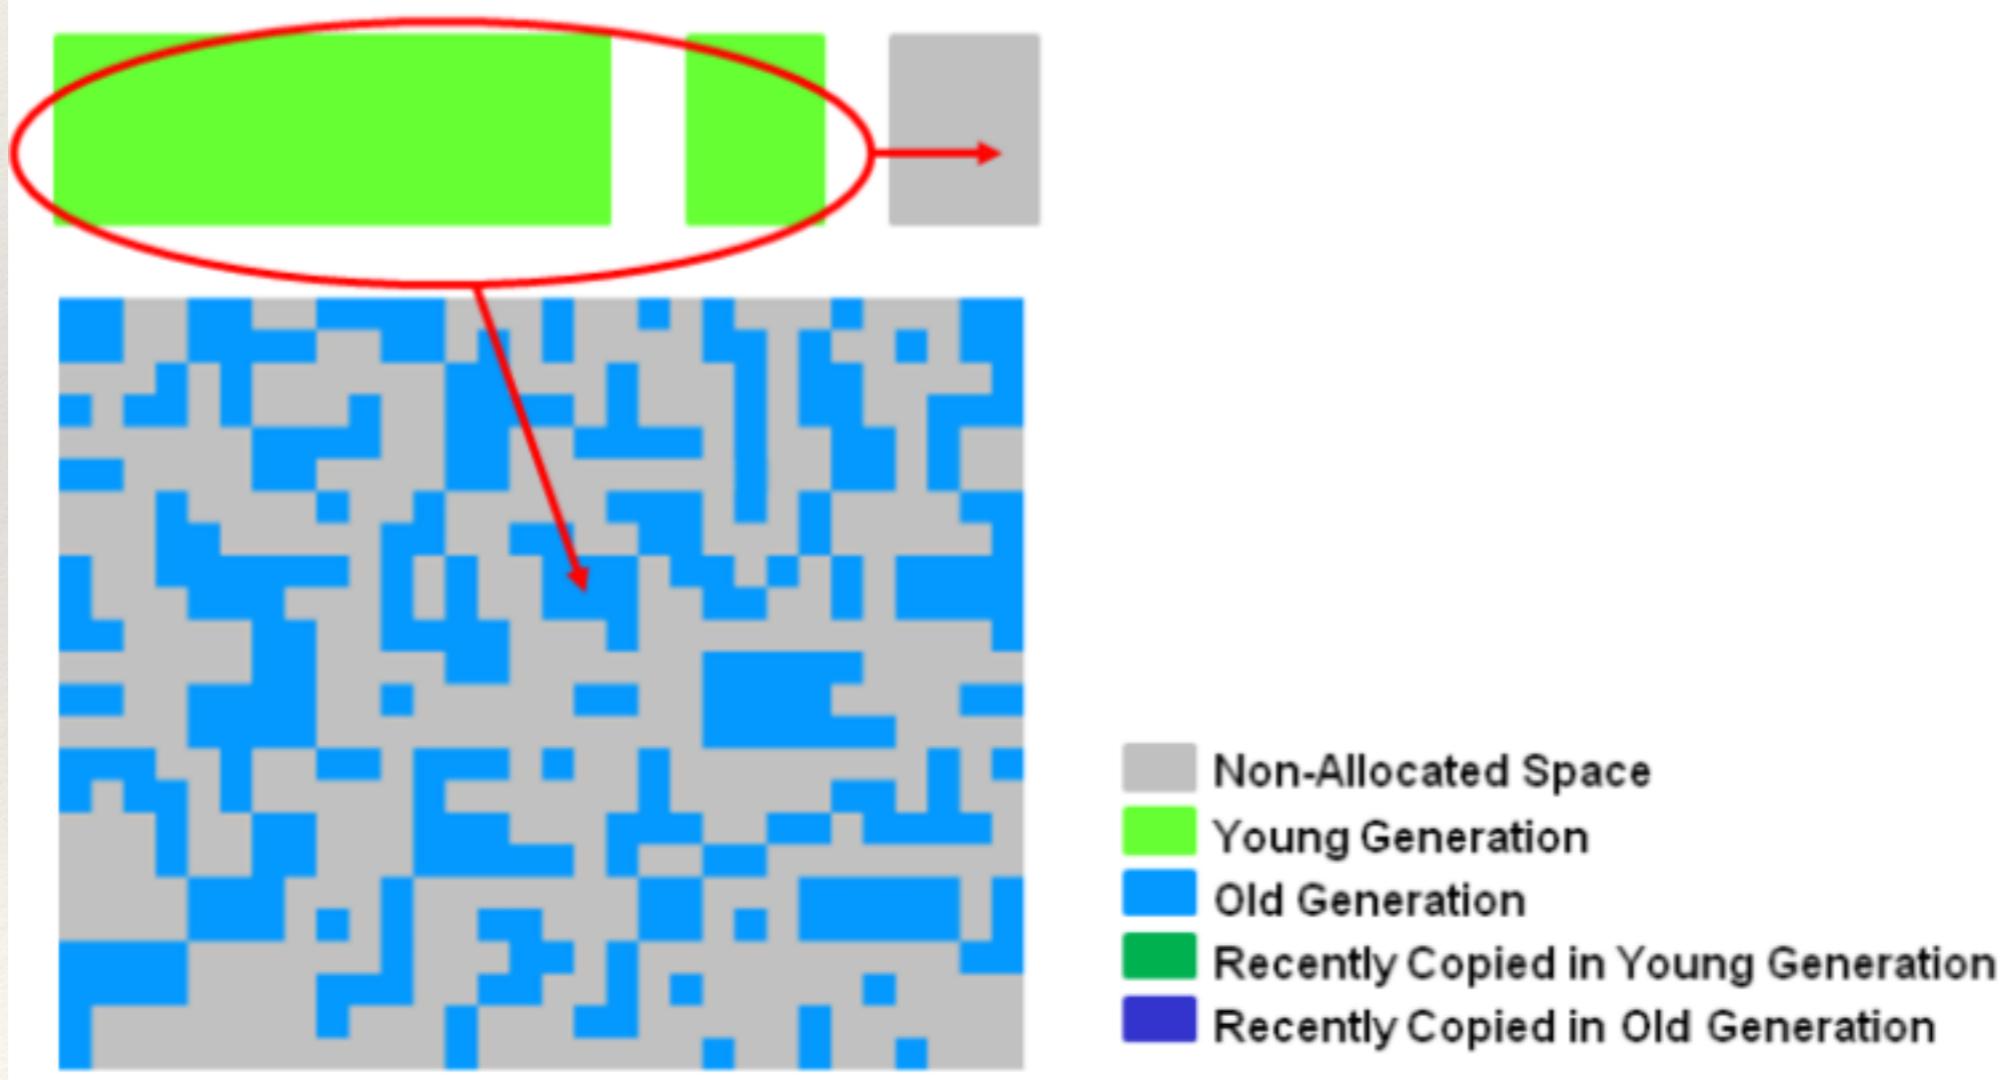
VM Options - CMS

How young GC Works



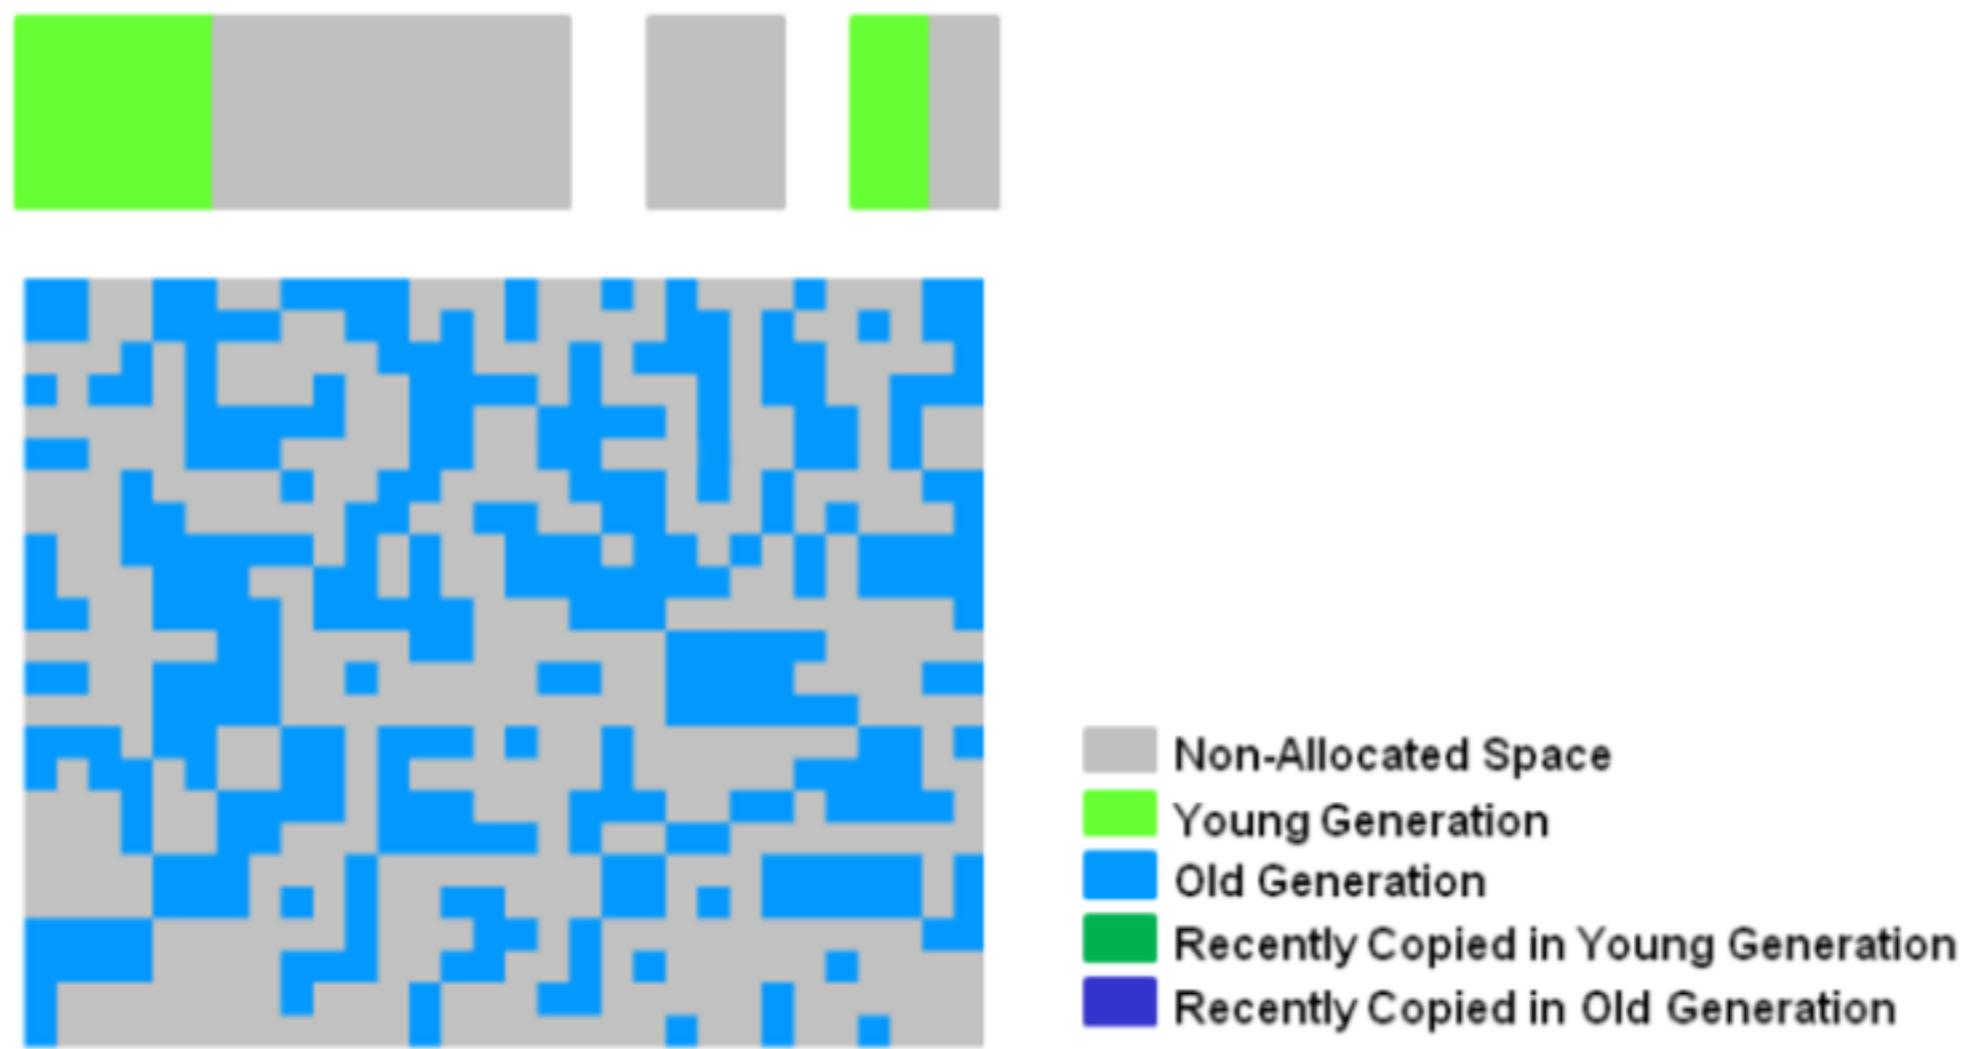
VM Options - CMS

Young Generation Collection



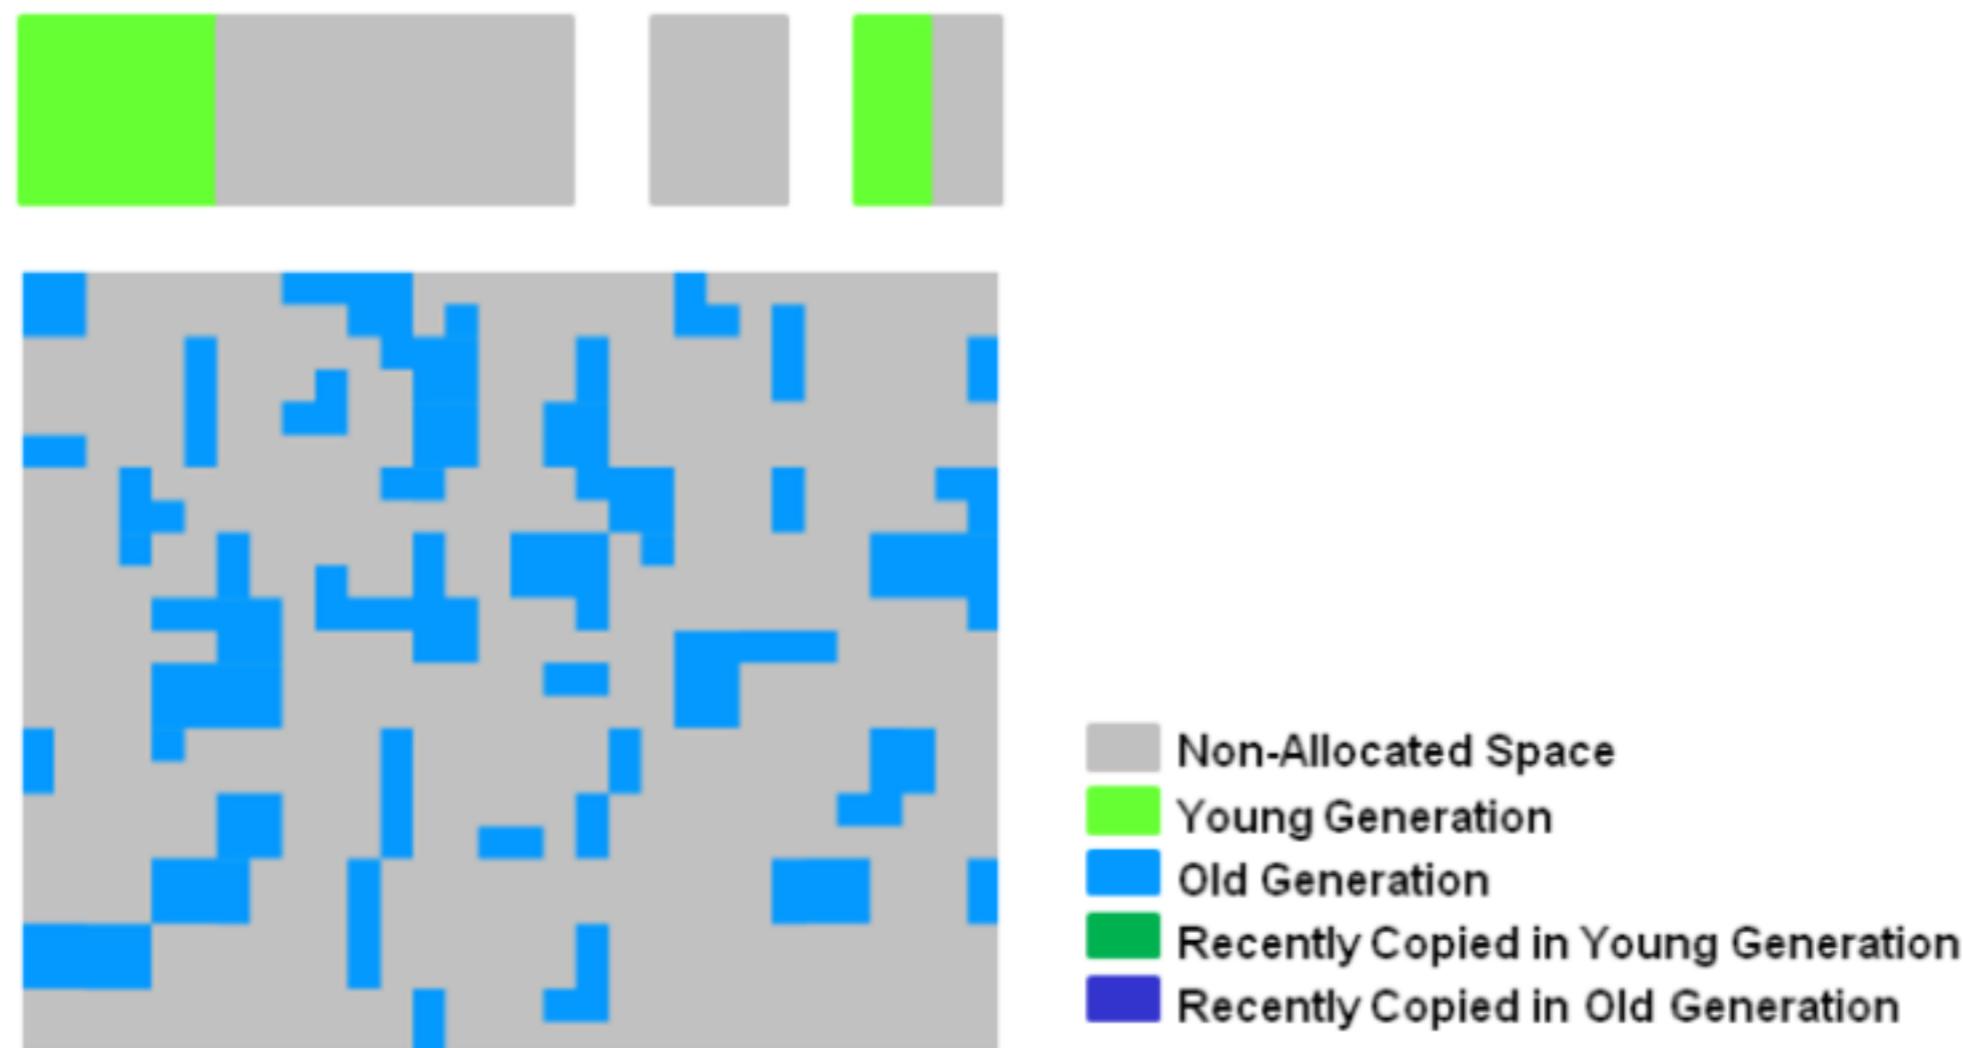
VM Options - CMS

Old Gen Collection – Concurrent Sweep



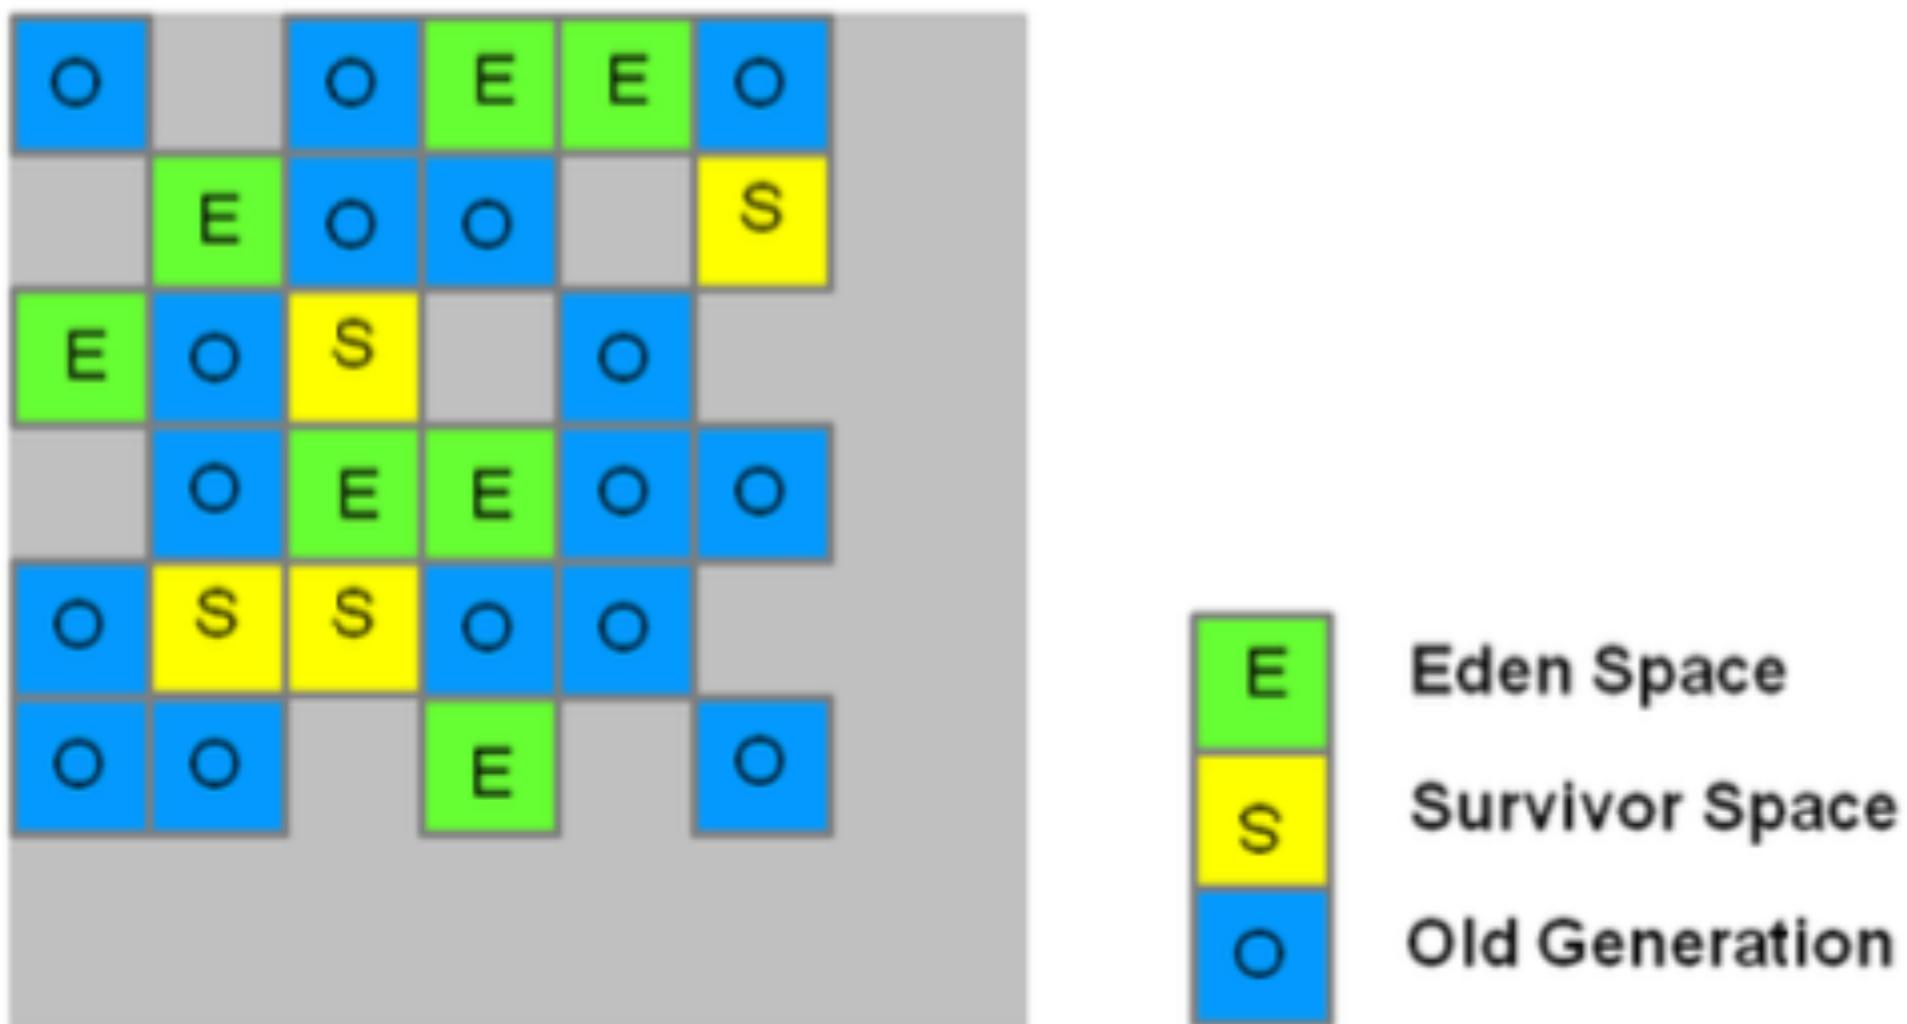
VM Options - CMS

Old Gen Collection – After Sweeping



VM Options - G1 GC

G1 Heap Allocation

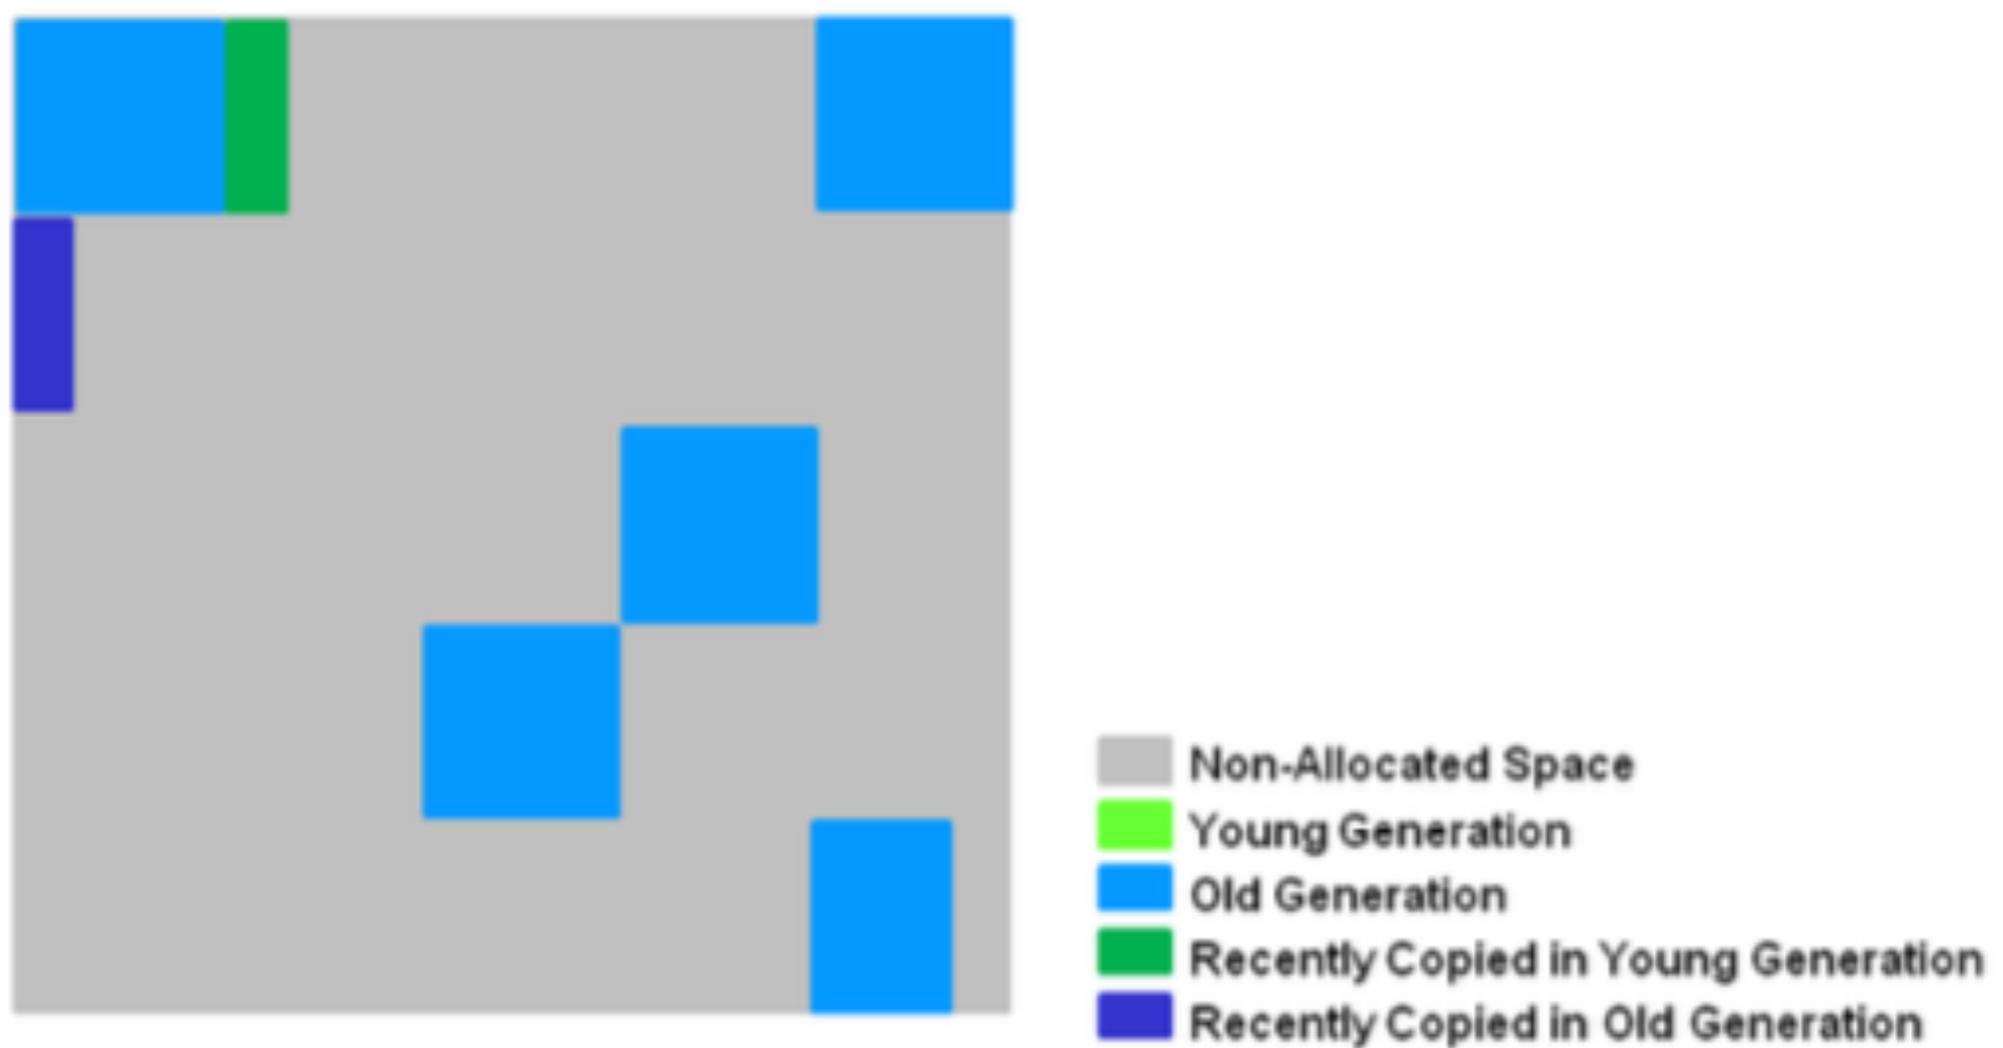


VM Options - G1 GC

Copying/Cleanup Phase



After Copying/Cleanup Phase



VM Options - G1 GC

- ❖ Based on the usage of the entire heap.
- ❖ `-XX:InitiatingHeapOccupancyPercent=n`
 - ❖ $n = 35$ # original value = 45
- ❖ `-XX:ParallelGCThreads=n`
 - ❖ if (cores > 8) cores * **5/8** else cores # From Oracle
- ❖ `-XX:ConcGCThreads=n,`
 - ❖ $n = \text{ParallelGCThreads} * \textcolor{red}{1/4}$ # From Oracle

G1 vs CMS

	G1	CMS
Heap compaction	Minor, Major, Full GC	Full GC
Heap Layout	Partition to regions	Young(eden, S0, S1) and Old
Heap Requirement	Normal Memory	Higher Memory

VM Options - JMX

- ❖ JMX = Java Management Extensions
- ❖ -Dcom.sun.management.jmxremote -
Dcom.sun.management.jmxremote.port={port} -
Dcom.sun.management.jmxremote.rmi.port={port} -
Dcom.sun.management.jmxremote.authenticate=false -
Dcom.sun.management.jmxremote.ssl=false -
Djava.rmi.server.hostname={ip or hostname}

VM Options - JMX

- ❖ VisualVM(jvisualvm)
- ❖ jmc

VM Options - Misc

- ❖ (A) `-XX:+UseCompressedStrings`(~~JDK 7, JDK 8~~)
- ❖ (B) `-XX:+AlwaysPreTouch`
- ❖ (C) `-XX:+UseStringCache`(~~JDK 8~~)
- ❖ (D) `-XX:+OptimizeStringConcat`
- ❖ Add B, C, and D will reduce around 10s from 6m 48s to 6m 37s in 48 cores and 256 GB RAM.

Summery

- ❖ The Java version was probably the simplest to implement, but **the hardest to analyze for performance.** Specifically the effects around **garbage collection were complicated and very hard to tune.** Since Scala runs on the JVM, it has the same issues. from [Loop Recognition in C++/Java/Go/Scala by Google](#)
- ❖ Knowledge of JVM(Java) still works for Scala.

Performance in production

- ❖ Performance in production = System Performance
 - ❖ System performance is the study of the entire system, **including all physical components and the full software stack.**
 - ❖ Performance engineering should ideally begin **before hardware is chosen or software is written.**
 - ❖ Performance is often **subjective.**
 - ❖ Performance can be a challenging discipline due to **the complexity of systems** and **the lack of a clear starting point for analysis.**
 - ❖ From Systems Performance: Enterprise and the Cloud by Brendan Gregg (Senior performance architect @ Netflix)

References

- ❖ <http://izquotes.com/>
- ❖ <http://www.azquotes.com/>
- ❖ <https://blog.codinghorror.com/how-to-become-a-better-programmer-by-not-programming/>
- ❖ <docs.scala-lang.org/overviews/collections/performance-characteristics.html>
- ❖ <https://www.gitbook.com/book/databricks/databricks-spark-knowledge-base/details>
- ❖ <https://0x0fff.com/spark-memory-management/>
- ❖ www.javaworld.com/article/2078635/enterprise-middleware/jvm-performance-optimization-part-2-compilers.html
- ❖ <https://github.com/dougqh/jvm-mechanics/blob/master/JVM%20Mechanics.pdf>
- ❖ [Java Performance by Charile Hunt and Binu John](#)
- ❖ [Loop Recognition in C++/Java/Go/Scala by Google](#)
- ❖ [Systems Performance: Enterprise and the Cloud by Brendan Gregg, Senior Performance Architect of Netflix](#)
- ❖ <https://databricks.com/blog/2015/05/28/tuning-java-garbage-collection-for-spark-applications.html>
- ❖ [Understanding Java Garbage Collection and what you can do about it by Gil Tene, CTO of Azul Systems](#)
- ❖ [Java Performance: The Definitive Guide by Scott Oaks, Architect of Oracle](#)
- ❖ [Martin Odersky: Scala with Style](#)
- ❖ [Scala Performance Considerations by Nermin Serifovic](#)
- ❖ [Scala for the Impatient by Cay Horstmann](#)
- ❖ [Parallel programming in Go and Scala A performance comparison by Carl Johnell](#)